arXiv:2611.08935v1 [cs.IR] 6 Nov 2026

SkinMind: Equipping Agentic RAG with Human-Inspired Graph Probing for Reliable Dermatology AI

Jing Hao[1,2], **Weinan Zhou**[1], **Zhihong Lu**[1], **Yaling Shen**[1,3], **Yiming Xu**[1,4], **Yefeng Liu**[1], **Siyuan Yan**[3], **Shuaibo Li**[6], **Ming Dai**[1,5], **Junjun He**[1], **Zuyi Zhu**[1], **Jinjie Gu**[1], **Lijun Wu**[1,*], **Jian Wang**[1,*]

[1]Ant Group, AFu Team, [2]The University of Hong Kong, [3]Monash University,
[4]University of Science and Technology of China, [5]Southeast University,
[6]The Hong Kong University of Science and Technology (Guangzhou)
*Corresponding authors.

Reliable dermatology AI requires access to verifiable evidence distributed across heterogeneous textual and visual sources. However, conventional retrieval-augmented generation (RAG) typically relies on one-shot retrieval from flat corpora, limiting its ability to adapt retrieval granularity, explore multi-hop relations, and integrate multimodal clinical evidence. We introduce **SkinMind**, a multimodal agentic RAG framework that equips policy models with human-inspired graph probing for autonomous evidence seeking. SkinMind organizes retrieval as an iterative interaction between an orchestrator and a graph environment. Through eight hierarchical graph-probing tools, the orchestrator progressively identifies relevant entities, inspects node-level information, traverses relations and paths, and retrieves source-text and visual evidence. This progressive-disclosure design allows the model to adaptively determine the retrieval modality, exploration direction, and evidence granularity while constructing a query-specific evidence subgraph. To support autonomous graph probing, we construct **SkinMind-KG** from a 62M-token corpus of dermatology guidelines and textbooks and 423K images, guided by a dermatologist-informed schema and a conflict-resolution mechanism. The resulting graph contains over 1.0M nodes and 2.0M edges, unifying textual, relational, and visual evidence for retrieval. Across six dermatology VQA benchmarks and two policy-model settings, SkinMind improves the average score over direct answering by 8.70 and 8.89 points with GPT-5.4-mini and Gemini-2.5-Flash, respectively, while consistently outperforming competing RAG methods. Together, SkinMind and SkinMind-KG provide an effective foundation for evidence-grounded multimodal reasoning in dermatology.

Date: November 6, 2026
Correspondence: Lijun Wu, lijun_wu@outlook.com; Jian Wang, bobblair.wj@antgroup.com
Keywords: Retrieval-augmented generation, agentic search, knowledge graphs, multimodal reasoning, dermatology

1 Introduction

Multimodal large language models (MLLMs) have demonstrated remarkable capabilities in language understanding, visual perception, and multi-step reasoning, showing considerable potential for addressing complex real-world tasks [28, 6, 12, 27, 33, 37, 1]. Dermatology represents a particularly promising application: many skin conditions are visually observable, and patients increasingly seek convenient access to professional medical guidance through online consultations. Accordingly, medical and dermatology-specific MLLMs have shown potential across tasks such as visual question answering, diagnostic decision-making, and treatment recommendation [10, 17, 16, 5, 31, 29, 44, 22]. Despite this promise, their deployment in safety-critical healthcare settings remains challenging. First, they may

generate fluent but factually incorrect statements, potentially leading to harmful clinical decisions. Second, their parametric knowledge is acquired from static training corpora and cannot be readily updated as clinical evidence and guidelines evolve. Third, their responses often lack verifiable references to supporting medical evidence, substantially limiting their reliability and interpretability.

Retrieval-augmented generation (RAG) offers a promising means of mitigating these limitations by grounding model responses in external, authoritative, and traceable medical evidence [9, 8, 13, 30, 35, 23, 36, 25, 24, 20, 3, 15]. Conventional RAG systems, however, typically perform a one-shot retrieval over a flat collection of textual chunks, presenting two limitations for complex medical reasoning [14]. First, one-shot retrieval relies solely on the initial query and cannot adapt to information needs that emerge during multi-step reasoning. Second, flat-chunk retrieval often provides fragmented evidence with insufficient relational context while introducing irrelevant information. Addressing the first limitation requires adaptive, multi-round retrieval, whereas addressing the second calls for a structured representation of external medical knowledge.

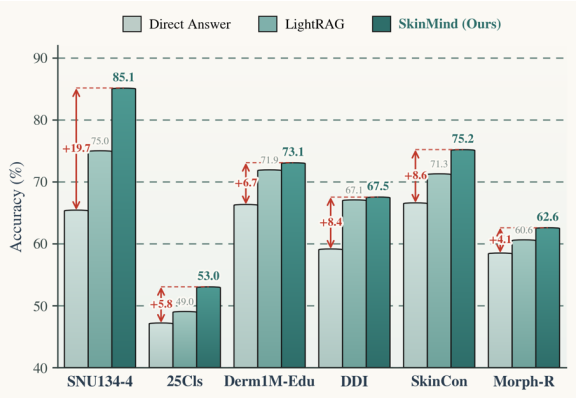

Figure 1: Performance comparison across six dermatology VQA benchmarks using Gemini-2.5-Flash as the policy model. SkinMind consistently outperforms direct answering and LightRAG.

Recent agentic RAG methods seek to address the first limitation by interleaving reasoning with iterative retrieval [8, 4]. Rather than relying on a fixed set of documents retrieved upfront, the model dynamically determines when additional evidence is needed and formulates new queries based on its intermediate reasoning state. Nevertheless, because these methods typically operate over isolated text chunks, they do not fully resolve the fragmentation and loss of relational context inherent in flat-text retrieval. Graph-based RAG improves upon flat-text retrieval by organizing knowledge into entities and relations [9, 13, 43]. However, many existing methods still treat the graph primarily as a retrieval index, selecting entities or subgraphs through embedding similarity or predefined expansion rules. Although the knowledge is graph-structured, the graph itself is rarely exposed to the model for interactive exploration. This contrasts with a human-inspired evidence-seeking process, in which a reasoner begins with a relevant concept, inspects its relations, follows informative paths, and adapts the exploration as new information emerges. This process follows the principle of *progressive disclosure* [32], whereby lightweight structural cues guide navigation, while detailed descriptions or source content are accessed only when needed. Compared with retrieving a fixed candidate set, such navigation allows the retrieval path, scope, and granularity to evolve with the reasoning process, facilitating the flexible construction of query-specific evidence subgraphs while reducing irrelevant context.

Inspired by this human-like evidence-seeking process, we propose **SkinMind**, a multimodal agentic RAG framework for autonomous and reliable dermatology assistance. Rather than treating the knowledge graph as a static retrieval index, SkinMind exposes it as an interactive environment and equips the agent with eight specialized graph-probing tools for interleaved reasoning and retrieval of textual, relational, and visual evidence. Given a clinical query, the agent autonomously identifies relevant entry nodes, inspects their relations, follows informative paths, and progressively accesses node descriptions, source content, and clinical images as needed. To provide a comprehensive and traceable foundation of domain-specific evidence, we construct **SkinMind-KG**, a large-scale heterogeneous dermatology knowledge graph containing over 1 million nodes and 2 million edges across 23 core entity types and 25 core relation types. Notably, SkinMind-KG explicitly links 301,561

images to disease entities, providing rich visual–conceptual associations for multimodal evidence seeking.

We conduct extensive experiments on six public dermatology VQA benchmarks spanning morphology recognition, disease diagnosis, and clinical reasoning in both closed- and open-ended settings. As illustrated in Fig. 1, SkinMind consistently outperforms direct answering across all six benchmarks, with gains of 4.1–19.7 points under Gemini-2.5-Flash. More broadly, across both policy-model settings, it surpasses LightRAG by an average of 5.75 points with GPT-5.4-mini and 3.60 points with Gemini-2.5-Flash. These results suggest that adaptive graph exploration can support more effective use of external knowledge and improve multimodal clinical reasoning. Our main contributions are summarized as follows:

- We propose **SkinMind**, a multimodal agentic RAG framework that treats a knowledge graph as an interactive environment, enabling interleaved reasoning and adaptive evidence seeking through eight graph-probing tools.
- We construct **SkinMind-KG**, a large-scale heterogeneous dermatology knowledge graph comprising over 1 million nodes and 2 million edges, with more than 300,000 images explicitly linked to disease entities.
- Extensive experiments on six dermatology VQA benchmarks across two policy-model settings demonstrate that SkinMind consistently outperforms direct answering and competing RAG methods, validating the effectiveness of adaptive graph exploration for evidence-grounded multimodal clinical reasoning.

2 Related Work

Retrieval-Augmented Generation. RAG augments language models with evidence from external corpora [21]. Adaptive methods interleave reasoning and retrieval, respond to generation uncertainty, incorporate self-critique, or route queries among retrieval strategies [34, 19, 3, 18]. GraphRAG, LightRAG, and HippoRAG organize evidence through entity graphs, community summaries, and associative retrieval [9, 13, 20]. Think-on-Graph 2.0 alternates graph and document retrieval, while Medical Graph RAG integrates medical sources and controlled vocabularies [24, 36]. Agentic approaches give models control over evidence acquisition: A-RAG exposes hierarchical passage-search interfaces, whereas RAGentA coordinates retrieval, answer generation, and verification across agents [8, 4]. SkinMind combines these directions by exposing a heterogeneous multimodal graph as an interactive environment. Its policy model selectively accesses nodes, relations, paths, source passages, and images at the granularity needed for each question.

Multimodal Models in Dermatology. SkinGPT-4 adapts a visual language model for interactive dermatology, while PanDerm provides domain-specific visual pretraining [44, 40]. Derm1M and MM-Skin supply large image–text resources for cross-modal representation, retrieval, and dermatology-specific vision–language modeling [39, 41]. DermoGPT emphasizes morphology-grounded diagnostic reasoning [29]. Tool-augmented systems complement these models with external memory and specialized components. SkinGPT-X coordinates specialist agents with an evolving case memory, while DermAgent combines perception tools, case and guideline retrieval, disease taxonomy, and evidence auditing [5, 22]. SkinMind uses a source-linked dermatology graph as its central retrieval environment, allowing the policy model to select the modality, relation, and evidence granularity during inference. This external knowledge can be updated independently of the policy model and traced to its supporting sources.

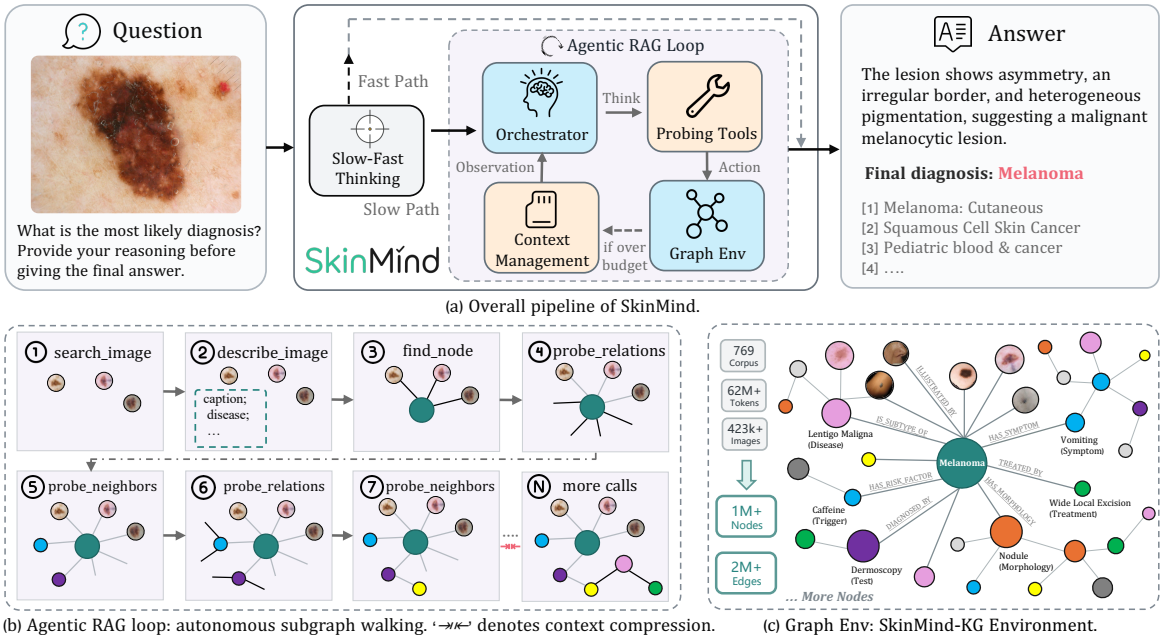

Figure 2: Overview of SkinMind. The orchestrator adaptively explores SkinMind-KG through hierarchical graph-probing tools to construct query-specific evidence subgraphs for response generation.

3 SkinMind

3.1 Overview

SkinMind is a multimodal agentic RAG framework equipped with a specialized suite of graph-probing tools for interleaved reasoning and evidence retrieval over a dermatology-specific graph environment. Given a clinical query, optionally accompanied by an image, SkinMind first routes the query through a slow–fast reasoning mechanism. Queries that can be answered reliably using the orchestrator's internal knowledge are handled by the fast path, whereas queries requiring external evidence, multimodal retrieval, or multi-step reasoning are directed to the slow path. The slow path implements an iterative Agentic RAG loop with four components: an *orchestrator*, a suite of *graph-probing tools*, a graph-based *environment*, and a *context manager*. The orchestrator acts as the central policy model: it reasons over the current context, selects and invokes appropriate tools, and incorporates the returned observations into subsequent decisions. Once sufficient evidence has been collected, it terminates the interaction loop and generates an evidence-grounded response for the user.

3.2 Slow–Fast Thinking

Let Q denote a textual clinical query and I an optional clinical image. A routing policy first determines whether external evidence seeking is necessary:

$$r = \pi_{\text{route}}(Q, I), \qquad r \in \{\textsc{Fast}, \textsc{Slow}\}. \tag{1}$$

The routing decision accounts for the user intent, query complexity, and the need for external or multimodal evidence. If $r = \textsc{Fast}$, the orchestrator directly generates a response. Otherwise, the query enters the graph-probing loop described below.

3.3 Agentic Graph-Probing Loop

Loop Formulation. For a query assigned to the slow path, the orchestrator alternates between reasoning, tool invocation, and answer generation based on the accumulated interaction history. At step t, the orchestrator state is defined as

$$s_t = (Q, I, \mathcal{H}_t)\,, \tag{2}$$

where Q is the textual query, I is an optional clinical image, and

$$\mathcal{H}_t = \big((a_0, o_0), \ldots, (a_{t-1}, o_{t-1})\big) \tag{3}$$

denotes the interaction history before step t. Here, a_i is the action taken at step i, and o_i is the corresponding observation returned by the graph environment.

Conditioned on s_t, the orchestrator either invokes a graph-probing tool or terminates the loop. For a non-terminal action, the selected tool operates on the graph environment and returns an observation, which is appended to the interaction history for the next reasoning step. The loop continues until the orchestrator selects FINISH or reaches the maximum interaction budget $L_{\max}$, producing a multi-turn trajectory

$$\begin{aligned}\tau = \big((s_0, a_0, o_0), (s_1, a_1, o_1), \ldots,\\ (s_{L-1}, a_{L-1}, o_{L-1}), (s_L, \textsc{Finish}, Y)\big),\end{aligned} \tag{4}$$

where $L \le L_{\max}$ and Y denotes the final response.

Tool Calling. Let $\mathcal{T} = \{T_1, \ldots, T_8\}$ denote the suite of graph-probing tools operating on textual and visual information in the graph environment. At each step, the orchestrator selects either a parameterized tool call

$$a_t = (T_{k_t}, \phi_t), \qquad T_{k_t} \in \mathcal{T}, \tag{5}$$

where ϕ_t denotes the tool arguments, or the terminal action FINISH. This decision is conditioned on the current state:

$$a_t \sim \pi_\theta(\cdot \mid s_t)\,. \tag{6}$$

The definitions and functionalities of these tools are detailed under *Hierarchical Graph-Probing Interfaces*.

Environment. Let $\mathcal{E}$ denote **SkinMind-KG**, our dermatology-specific multimodal graph environment. For each non-terminal action, the environment executes the corresponding tool call and returns an observation:

$$o_t = \mathcal{E}(a_t), \qquad a_t \neq \textsc{Finish}. \tag{7}$$

The construction and organization of SkinMind-KG are detailed under *SkinMind-KG Construction and Statistics*.

Context Management. To control context growth during multi-step graph exploration, SkinMind monitors both individual tool observations and the accumulated context. If a single observation exceeds 4,096 tokens, the orchestrator compresses it before incorporating it into the interaction history. If the accumulated context exceeds 16,384 tokens, the orchestrator summarizes earlier interactions to reduce the overall context length. These operations preserve query-relevant evidence and retrieval provenance required for subsequent reasoning and graph exploration.

3.4 Hierarchical Graph-Probing Interfaces

SkinMind equips the orchestrator with eight graph-probing tools for retrieving and exploring textual, relational, and visual evidence in SkinMind-KG. As summarized in Table 1, six tools support textual retrieval and graph exploration, while two tools support visual retrieval and inspection. Following the principle of *progressive disclosure* [32, 8], the interfaces initially return compact retrieval results and expose more detailed attributes, relations, and source evidence only when requested, thereby reducing unnecessary context overhead.

Table 1: **Overview of SkinMind's graph-probing tool suite.** The first six tools support textual retrieval and relation-aware graph exploration, while the last two support visual retrieval and metadata inspection.

Tool	Function	Arguments	Output
`find_node`	Retrieve graph nodes using a textual query	Query, top-K	Node IDs, names, and types
`read_node`	Retrieve the structured information of a node	Node ID	Node attributes and description
`read_chunks`	Retrieve source-text chunks associated with a node	Node ID, chunk limit, character limit	Source-text chunks
`probe_relations`	List relation types associated with a node	Node ID	Relation types and neighbor counts
`probe_neighbors`	Expand a node along a specified relation	Node ID, relation, top-K	Neighbor IDs, names, and types
`probe_path`	Search for relational paths between two nodes	Source ID, target ID, hop limit	Node and relation sequences
`search_image`	Retrieve image nodes using a text or image query	Text/image query, top-K	Image-node IDs and metadata
`describe_image`	Retrieve the metadata of an image node	Image-node ID	Caption, morphology, and disease labels

The orchestrator can enter the graph through textual node retrieval with `find_node` or through text-to-image or image-to-image retrieval with `search_image`. For image-to-image retrieval, `search_image` encodes the query image with DermLIP [39] and searches the image-node index by cosine similarity, returning image identifiers, similarity scores, and source types. Depending on its information needs, the orchestrator may inspect captions, clinical attributes, and aligned disease labels using `describe_image`. The image-to-disease associations established during graph construction allow these labels to serve as anchors for locating `Disease` nodes through `find_node`.

To gather further evidence, the orchestrator can inspect node details using `read_node`, explore graph structure using `probe_relations`, `probe_neighbors`, and `probe_path`, and access fine-grained source evidence using `read_chunks`. Through iterative tool calls, it adaptively determines the retrieval modality, exploration direction, and level of evidence granularity based on its evolving information needs. The retrieved nodes, relations, source chunks, and image records collectively form a query-specific evidence subgraph for subsequent reasoning and response generation. Figure 2(b) illustrates an example graph-probing trajectory.

4 SkinMind-KG Construction and Statistics

4.1 Knowledge Graph Construction

We construct the textual component of SkinMind-KG from a large-scale dermatology corpus comprising 227 clinical guidelines and 542 textbooks published since 2015. We collected dermatology guidelines and textbooks from international online sources, selecting the latest available guideline versions and textbook editions at the time of collection. We parse the raw documents using MinerU [26]

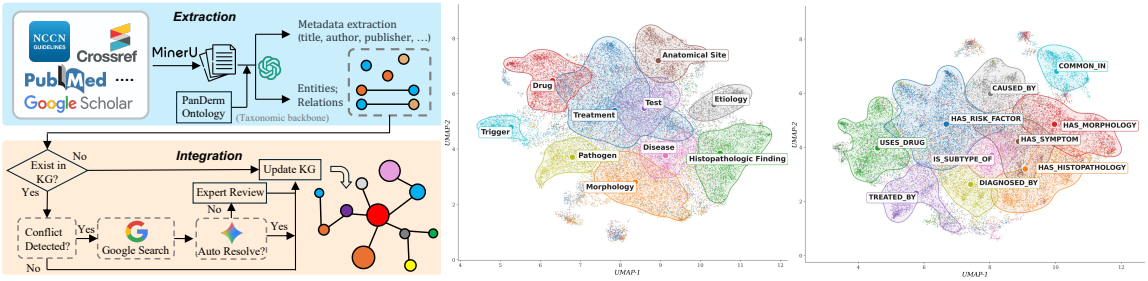

(a) Construction pipeline of the textual SkinMind-KG. (b) UMAP of entity embeddings by semantic theme. (c) UMAP of relation embeddings by semantic theme.

Figure 3: Construction and semantic organization of SkinMind-KG. The ontology-guided pipeline integrates knowledge with conflict resolution, while UMAP projections of entity and relation embeddings reveal clusters across semantic categories.

and employ a lightweight LLM to extract document-level metadata, including titles, authors, publishers, publication years, and editions. To establish an authoritative disease taxonomy, we adopt the PanDerm ontology [40], which comprises 127 disease concepts and their curated hierarchical relationships. We encode these relationships as `IS_SUBTYPE_OF` relations and preload the resulting hierarchy into SkinMind-KG as its initial taxonomic backbone. Through iterative consultations with senior dermatologists, we further define a schema containing 23 entity types and 25 relation types to guide information extraction. Following LightRAG [13], we segment the approximately 62M token corpus into chunks with a length of 1,200 tokens. GPT-5.4-mini is then used to extract entities, relations, and their descriptions from each chunk according to the predefined schema. During corpus extraction, the initial disease hierarchy can be extended with additional disease entities and subtype relations.

Entity and relation extraction is restricted to statements explicitly supported by the source text. The source context guides the interpretation of clinical qualifiers and the polarity of these statements. Extracted entities and textual relations retain links to their supporting source chunks, which preserve document-level provenance. Through these links, the orchestrator can retrieve supporting passages and their document identifiers using `read_chunks` when node-level descriptions are insufficient, providing access to the original clinical context.

Conflict Detection and Resolution. Integrating heterogeneous sources may introduce contradictory claims about the same entity or relation. We therefore incorporate a *detect-then-resolve* procedure into the construction pipeline. A conflict is identified only when two claims make contradictory assertions about the same fact under the same clinical context. Claims associated with different clinical contexts are retained as context-specific evidence rather than treated as conflicts. For each detected conflict, Gemini-3.1-Pro [12] formulates targeted queries and retrieves relevant evidence through Google Search. It then evaluates the retrieved evidence to determine whether the conflict can be reliably resolved. Resolvable cases are automatically reconciled, whereas ambiguous cases are referred to domain experts. In total, the procedure processed 1,375 description-level conflicts, automatically resolving 1,366 cases (99.3%) and referring the remaining nine for expert review. The overall procedure is illustrated in Figure 3(a).

Visual Modality Integration. After constructing the textual graph, we augment SkinMind-KG with 423,016 image nodes from Derm1M [39], most of which are accompanied by disease labels. To integrate these images into the graph, we align their disease labels with textual graph nodes of type `Disease` through a three-stage procedure. First, normalized image labels are matched exactly against the names of disease nodes. Second, for unmatched labels, Qwen3-Embedding-4B [42] is used to retrieve the top-10 semantically closest disease nodes. Third, GPT-5.4-mini disambiguates among these candidates and determines whether a reliable label-to-node alignment can be established. Based on the resulting

alignments, image nodes are connected to the corresponding disease nodes through ILLUSTRATED_BY relations. Images without a reliable alignment remain unlinked to disease nodes. To support multimodal retrieval, the images and their associated textual labels are encoded using the visual and textual branches of DermLIP [39], respectively. The resulting representations are indexed to support text-to-image and image-to-image retrieval over the image nodes.

4.2 Graph Statistics and Structural Characteristics

The resulting knowledge graph, **SkinMind-KG**, is a multimodal, heterogeneous graph containing 1,006,079 nodes and 2,038,419 edges. The 23 core entity types account for 99% of all non-image nodes, while the 25 core relation types account for 99% of all edges excluding ILLUSTRATED_BY. The graph includes 423,016 image nodes, of which 301,561 are connected to disease nodes through ILLUSTRATED_BY relations, providing explicit associations between clinical images and dermatological concepts. SkinMind-KG is unweighted and undirected. Its largest connected component contains 80.0% of all nodes and 99.5% of all edges, indicating that nearly all relational evidence is organized within a shared connected subgraph. Table 2 reports the node and edge counts, degree distribution, and connected-component statistics. The average degree is 4.05; among non-isolated nodes, the median and 99th-percentile degrees are 1 and 47, respectively, while the maximum degree is 64,345. This variation in neighborhood size motivates relation-level probing to guide selective neighbor expansion, enabling the orchestrator to adapt the breadth of exploration to the local graph structure.

Table 2: Node and edge counts, degree distribution, and connected-component statistics of SkinMind-KG after multimodal integration.

Property	**Value**
Scale & Composition	
Nodes	1,006,079
Edges	2,038,419
Graph type	Undirected, unweighted
99%-coverage entity types	23
99%-coverage relation types	25
Image nodes	423,016 (42.0%)
Linked to disease nodes	301,561 (71.3%)
Isolated image nodes	121,455 (28.7%)
Degree & Density	
Average degree	4.05
Median / p90 / p99 degree	1 / 5 / 47
Non-isolated nodes	819,864 (81.5%)
Maximum degree	64,345
Degree-one nodes	432,054 (52.7%)
Graph density	4.0×10^{-6}
Connectivity	
Connected components	191,613
Isolated nodes	186,215 (18.5%)
LCC (nodes)	805,181 (80.0%)
LCC (edges)	2,029,050 (99.5%)

Note: LCC denotes the largest connected component. Degree percentiles and the proportion of degree-one nodes use non-isolated nodes. Linked and isolated image percentages use all image nodes. Core type coverage excludes image nodes and ILLUSTRATED_BY edges, respectively.

5 Experiments

5.1 Experimental Setting

Datasets. We evaluate SkinMind on six public dermatology VQA benchmarks spanning both closed- and open-ended settings. SkinCon [29], 25Cls [29], Derm-Edu [29], DDI [7], and SNU134-4 [29] are formulated as multiple-choice benchmarks, whereas Morph-R [29] requires free-form responses. Collectively, these benchmarks assess morphology recognition, disease diagnosis, and clinical reasoning across different response formats. We use SkinMind-KG as the retrieval corpus. To prevent benchmark leakage from image overlap, we first remove exact duplicates based on MD5 hashes. We then exclude retrieved image nodes whose DermLIP cosine similarity to the query image exceeds 0.98, thereby filtering potential near duplicates. This two-stage procedure is applied identically to all RAG methods to ensure a fair comparison.

Baselines and Evaluation. We organize the baselines into three groups. (i) *Standard baselines* comprise zero-shot inference without retrieval (*Direct Answer*) and conventional single-round retrieval augmen-

Table 3: Comparison across six dermatology VQA benchmarks (higher is better). Standalone models serve as contextual references and are not directly comparable with the shared SkinMind protocol. Bold and underline mark the best and second-best scores in each column within each of the three blocks. Δ denotes SkinMind's absolute gain over the best other method in the same policy block, selected independently per column and including Direct Answer. Average is the unweighted mean across six benchmarks; its Δ uses the highest baseline Average.

Setting	Model / Method	SkinCon	25Cls	Derm-Edu	DDI	SNU134-4	Morph-R	Average
Strong Reference Models								
Standalone	GPT-5.5 [28]	72.01	50.40	**72.48**	**65.79**	**81.50**	**66.30**	**68.08**
	Claude-Sonnet-4.5 [2]	59.20	34.15	53.64	52.90	68.75	54.37	53.84
	DermoGPT-RL [29]	**76.67**	**64.60**	69.68	62.80	60.00	65.48	66.54
	SkinVL-PubMM [41]	61.31	38.75	51.12	48.93	54.58	54.62	51.55
	Lingshu-32B [38]	52.39	38.40	30.29	34.91	45.83	49.55	41.90
GPT-5.4-mini as Policy Model								
Standard	Direct Answer	63.67	42.22	63.42	50.30	73.33	62.24	59.20
	Naive RAG [11]	67.25	50.40	66.79	52.90	71.25	62.50	61.85
Graph-based RAG	GraphRAG [9]	64.24	46.55	65.80	54.12	71.25	60.28	60.37
	LightRAG [13]	64.95	45.80	68.85	58.38	72.50	62.39	62.15
Agentic RAG	A-RAG [8]	66.16	49.50	63.52	53.81	70.00	62.33	60.89
	RAGentA [4]	58.63	43.85	65.58	49.70	65.83	58.71	57.05
	SkinMind	**68.84**	**54.05**	**70.60**	**65.09**	**85.00**	**63.84**	**67.90**
	Δ vs. Best Baseline	**+1.59**	**+3.65**	**+1.75**	**+6.71**	**+11.67**	**+1.34**	**+5.75**
Gemini-2.5-Flash as Policy Model								
Standard	Direct Answer	66.59	47.20	66.33	59.15	65.42	58.49	60.53
	Naive RAG [11]	64.93	49.75	70.46	60.67	74.58	61.89	63.71
Graph-based RAG	GraphRAG [9]	59.64	47.65	69.79	61.13	71.67	59.71	61.60
	LightRAG [13]	71.28	49.05	71.92	67.07	75.00	60.62	65.82
Agentic RAG	A-RAG [8]	70.54	46.85	72.56	62.80	72.50	62.14	64.57
	RAGentA [4]	70.58	49.10	72.53	62.80	76.67	60.97	65.44
	SkinMind	**75.18**	**53.03**	**73.06**	**67.53**	**85.11**	**62.58**	**69.42**
	Δ vs. Best Baseline	**+3.90**	**+3.28**	**+0.50**	**+0.46**	**+8.44**	**+0.44**	**+3.60**

tation (*Naive RAG*) [11]. (ii) *Graph-based RAG methods* include GraphRAG [9] and LightRAG [13], which leverage graph-structured knowledge for retrieval. (iii) *Agentic RAG methods* include A-RAG [8] and RAGentA [4], which interleave model reasoning with iterative retrieval. To ensure a fair comparison, all methods use the same answering model and inference settings and receive identical task instructions and output-format requirements. All retrieval-augmented methods operate on the same SkinMind-KG through a shared retrieval backend. Prompt differences are limited to method-specific retrieval and tool-use instructions. Following the official evaluation protocol of DermoGPT-RL [29], we report accuracy on the five multiple-choice benchmarks and evaluate open-ended responses on Morph-R using Gemini-2.5-Pro under an LLM-as-a-judge protocol. We use the official DermoGPT evaluation code without modification.[1] The judge compares each response with its reference answer without access to the input image. It assesses clinical reasoning, final diagnosis, and closed-set morphology output, which the official evaluator combines into the reported Morph-R score.

Implementation Details. We instantiate each method with two policy models: GPT-5.4-mini [27] and Gemini-2.5-Flash [6]. Non-agentic methods retrieve the top five textual items and the top five image items, whereas agentic methods allow the policy model to determine k dynamically at each retrieval

[1] https://github.com/mendicant04/DermoGPT

round, with $k = 5$ as the default. Textual and visual retrieval is performed using Qwen3-Embedding-4B [42] and DermLIP [39], respectively, based on cosine similarity. Agentic methods are limited to 20 retrieval rounds per question. For generation, we set the temperature to 1.0 and the maximum context length to 32,768 tokens. All methods instantiated with the same policy model use identical generation settings. The topology of SkinMind-KG is stored in GraphML format, whereas Qdrant stores embedding indices for the shared retrieval backend.

5.2 Main Results

Table 3 presents the main experimental results across six benchmarks. We highlight three key observations.

SkinMind delivers consistent performance gains. SkinMind consistently outperforms the corresponding *Direct Answer* baseline. With GPT-5.4-mini, it improves the average score from 59.20 to 67.90, a gain of 8.70 points; with Gemini-2.5-Flash, it increases the score from 60.53 to 69.42, a gain of 8.89 points. Most retrieval-augmented methods also improve over direct answering, supporting the value of SkinMind-KG as an external knowledge source. Notably, SkinMind outperforms the competing RAG methods under both policy-model settings, indicating that adaptive graph probing enables more effective use of this knowledge.

Adaptive graph probing enables more effective retrieval. Compared with LightRAG, SkinMind improves the average score by 5.75 points with GPT-5.4-mini and 3.60 points with Gemini-2.5-Flash. Meanwhile, the gains of other RAG methods vary across benchmarks and policy models, suggesting that retrieval effectiveness depends not only on access to external knowledge but also on how that knowledge is explored. SkinMind adaptively selects entry nodes, probes relations, and traverses paths to gather query-specific textual and visual evidence at the required granularity.

SkinMind is competitive with strong standalone reference models. Although not directly comparable, standalone models provide useful reference points. SkinMind with GPT-5.4-mini achieves an average score of 67.90, comparable to GPT-5.5's 68.08, while outperforming DermoGPT-RL, SkinVL-PubMM, and Lingshu-32B by 1.36, 16.35, and 26.00 points, respectively. This result highlights the effectiveness of graph-based evidence seeking in narrowing the gap between the policy model and stronger standalone models.

5.3 Ablation on Graph Probing Tools

To examine the roles of graph exploration, image retrieval, and progressive disclosure, we compare SkinMind with three ablation variants on DDI and SNU134-4 using GPT-5.4-mini. *w/o Graph Exploration* disables `probe_relations`, `probe_neighbors`, and `probe_path`, while preserving textual and visual node retrieval. *w/o Image Retrieval* disables `search_image` and `describe_image`. In *w/o Progressive Disclosure*, `read_node` directly returns the associated source chunks normally retrieved through a separate `read_chunks` call.

Table 4: Ablations with GPT-5.4-mini on DDI and SNU134-4.

Method	DDI	SNU134-4
SkinMind (Full)	**65.09**	**85.00**
w/o Graph Exploration	62.65 (−2.44)	82.50 (−2.50)
w/o Image Retrieval	63.26 (−1.83)	79.17 (−5.83)
w/o Progressive Disclosure	64.48 (−0.61)	80.83 (−4.17)

Variants without graph exploration, image retrieval, or progressive disclosure show lower accuracy on both benchmarks (Table 4). On DDI, *w/o Graph Exploration* shows the largest gap at 2.44 points. On SNU134-4, *w/o Image Retrieval* and *w/o Progressive Disclosure* show more pronounced gaps of 5.83

and 4.17 points, respectively. These patterns are consistent with complementary roles for relational context, visual references, and selective evidence disclosure in diagnostic reasoning. Together, these capabilities allow the orchestrator to adapt both the evidence source and the retrieval depth to the diagnostic query.

5.4 Illustrative Case Study

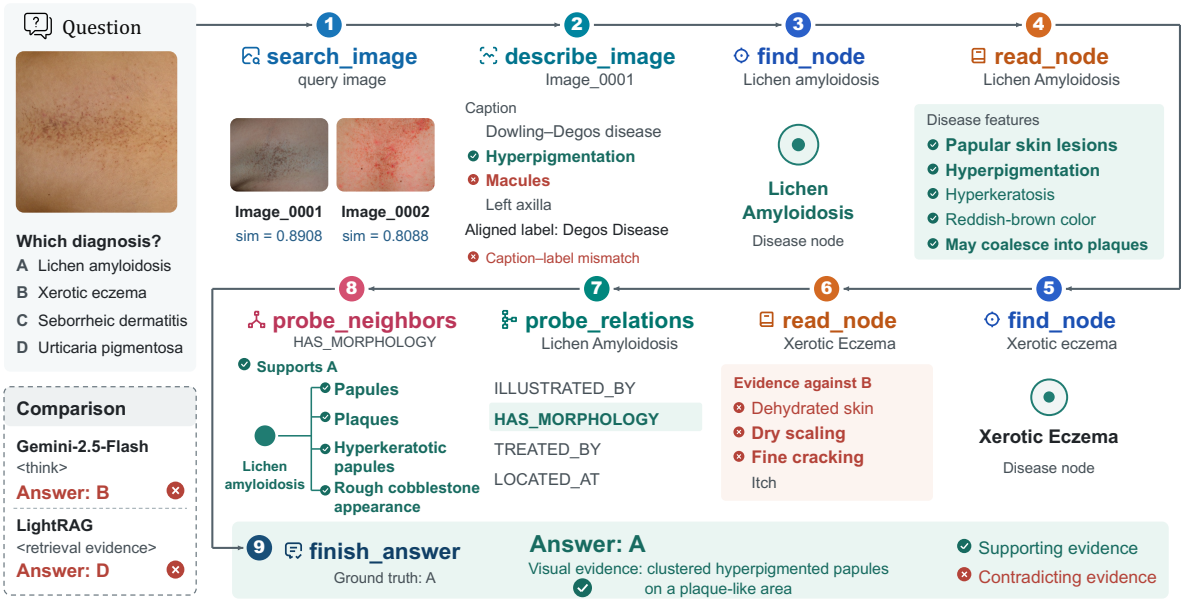

Figure 4: An illustrative graph-probing trajectory of SkinMind. By iteratively retrieving visual and textual evidence and comparing plausible diagnoses, SkinMind correctly identifies lichen amyloidosis, whereas Gemini-2.5-Flash and LightRAG produce incorrect predictions.

Figure 4 illustrates a reasoning trajectory of SkinMind for dermatological diagnosis. Given a clinical image and a multiple-choice question, SkinMind progressively retrieves visually similar cases, extracts dermatological concepts, and explores relevant disease nodes to gather supporting evidence and eliminate competing diagnoses. By integrating the observed reddish-brown hyperpigmented papules and macules, coalescing lesions, and signs of chronic scratching with the retrieved knowledge, SkinMind correctly identifies the condition as lichen amyloidosis. In contrast, Gemini-2.5-Flash and LightRAG select incorrect answers. This example illustrates how SkinMind's iterative evidence-seeking process supports fine-grained, evidence-grounded clinical reasoning.

5.5 Context Efficiency Analysis

Figure 5 compares the average token consumption of SkinMind against five representative retrieval-augmented baselines on the Morph-R benchmark. The stacked bars show input and output tokens (left axis), while the line reports the average number of retrieved evidence tokens in the model context (right axis).

Evidence efficiency. Most notably, SkinMind achieves the strongest task performance while using the fewest evidence tokens among all evaluated methods, averaging 1,380 tokens per question. By comparison, LightRAG uses 7,901 evidence tokens, 5.7× as many as SkinMind. This pattern suggests that SkinMind's gains are associated with more targeted evidence selection rather than a larger retrieved context. Its hierarchical graph-probing interfaces support selective access to relevant textual and visual evidence at different levels of granularity.

Reasoning-oriented token allocation. SkinMind allocates 21% of its token usage to model outputs, amounting to 5,576 tokens per question. These outputs include both the final answer and the intermediate decisions involved in navigating the graph and selecting evidence. In comparison, LightRAG allocates 2% of its token usage (291 tokens) to model outputs, with 11,802 input tokens consumed through upfront context construction. Aggregated across calls, total usage is approximately 26.1k tokens for SkinMind and 12.1k for LightRAG. SkinMind's multi-round inputs include the interaction history and tool instructions alongside retrieved evidence. This allocation characterizes its iterative evidence-acquisition process: the model progressively selects what to retrieve and uses the resulting observations to guide subsequent graph-probing actions.

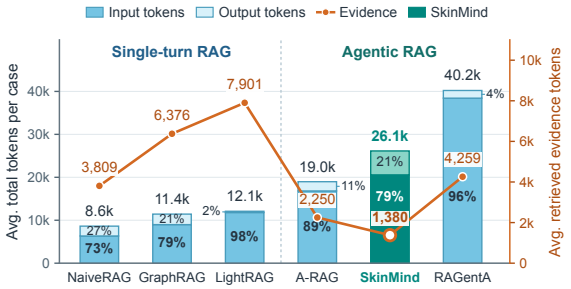

Figure 5: Token-efficiency comparison across RAG methods on Morph-R using Gemini-2.5-Flash as the policy model. Stacked bars show average input and output tokens, while the line indicates average retrieved evidence tokens. SkinMind retrieves the most compact evidence set.

6 Conclusion

We introduced **SkinMind**, a multimodal agentic RAG framework that interactively probes a dermatology knowledge graph for textual, relational, and visual evidence. Equipped with hierarchical graph-probing tools, SkinMind interleaves reasoning with adaptive exploration of textual, relational, and visual evidence, progressively constructing query-specific evidence subgraphs. We also developed **SkinMind-KG** to provide a comprehensive and traceable foundation of domain-specific knowledge. Across six dermatology VQA benchmarks and two policy models, SkinMind consistently outperforms direct answering and competing RAG methods while retrieving more targeted evidence. These results support adaptive graph probing as a promising approach to evidence-grounded dermatology reasoning.

References

[1] Marah Abdin, Jyoti Aneja, Harkirat Behl, Sébastien Bubeck, Ronen Eldan, Suriya Gunasekar, Michael Harrison, Russell J Hewett, Mojan Javaheripi, Piero Kauffmann, et al. Phi-4 technical report. *arXiv preprint arXiv:2412.08905*, 2024.

[2] Anthropic. Introducing Claude Sonnet 4.5. https://www.anthropic.com/news/claude-sonnet-4-5, 2025.

[3] Akari Asai, Zeqiu Wu, Yizhong Wang, Avirup Sil, and Hannaneh Hajishirzi. Self-RAG: Learning to retrieve, generate, and critique through self-reflection. In *The Twelfth International Conference on Learning Representations*, 2024. URL https://openreview.net/forum?id=hSyW5go0v8.

[4] Ines Besrour, Jingbo He, Tobias Schreieder, and Michael Färber. Ragenta: Multi-agent retrieval-augmented generation for attributed question answering. *arXiv preprint arXiv:2506.16988*, 2025.

[5] Zhangtianyi Chen, Yuhao Shen, Florensia Widjaja, Yan Xu, Liyuan Sun, Zijian Wang, Hongyi Chen, Wufei Dai, and Juexiao Zhou. Skingpt-x: A self-evolving collaborative multi-agent system for transparent and trustworthy dermatological diagnosis. *arXiv preprint arXiv:2603.26122*, 2026.

[6] Gheorghe Comanici, Eric Bieber, Mike Schaekermann, et al. Gemini 2.5: Pushing the frontier with advanced reasoning, multimodality, long context, and next generation agentic capabilities, 2025. URL https://arxiv.org/abs/2507.06261.

[7] Roxana Daneshjou, Kailas Vodrahalli, Roberto A Novoa, Melissa Jenkins, Weixin Liang, Veronica Rotem-

berg, Justin Ko, Susan M Swetter, Elizabeth E Bailey, Olivier Gevaert, et al. Disparities in dermatology ai performance on a diverse, curated clinical image set. *Science advances*, 8(31):eabq6147, 2022.

[8] Mingxuan Du, Benfeng Xu, Chiwei Zhu, Shaohan Wang, Pengyu Wang, Xiaorui Wang, and Zhendong Mao. A-rag: Scaling agentic retrieval-augmented generation via hierarchical retrieval interfaces. *arXiv preprint arXiv:2602.03442*, 2026.

[9] Darren Edge, Ha Trinh, Newman Cheng, Joshua Bradley, Alex Chao, Apurva Mody, Steven Truitt, Dasha Metropolitansky, Robert Osazuwa Ness, and Jonathan Larson. From local to global: A graph RAG approach to query-focused summarization, 2024. URL https://arxiv.org/abs/2404.16130.

[10] Yuxuan Fan, Jing Hao, Hong Chen, Jiahao Bao, Yihua Shao, Yuci Liang, Kuo Feng Hung, and Hao Tang. Oralgpt-plus: Learning to use visual tools via reinforcement learning for panoramic x-ray analysis. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, pages 35373–35383, 2026.

[11] Yunfan Gao, Yun Xiong, Xinyu Gao, Kangxiang Jia, Jinliu Pan, Yuxi Bi, Yi Dai, Jiawei Sun, Meng Wang, and Haofen Wang. Retrieval-augmented generation for large language models: A survey. *arXiv preprint arXiv:2312.10997*, 2023.

[12] Google DeepMind. Gemini 3.1 Pro model card. https://deepmind.google/models/model-cards/gemini-3-1-pro/, 2026.

[13] Zirui Guo, Lianghao Xia, Yanhua Yu, Tu Ao, and Chao Huang. Lightrag: Simple and fast retrieval-augmented generation. *arXiv preprint arXiv:2410.05779*, 2024.

[14] Haoyu Han, Li Ma, Yu Wang, Harry Shomer, Yongjia Lei, Zhisheng Qi, Kai Guo, Zhigang Hua, Bo Long, Hui Liu, et al. Rag vs. graphrag: A systematic evaluation and key insights. *arXiv preprint arXiv:2502.11371*, 2025.

[15] Jing Hao, Siyuan Dai, Yongxin Zhang, Yuci Liang, Jiamin Wu, Jiahao Bao, Yuxuan Fan, Zanting Ye, Yanpeng Sun, Xinyu Zhang, Ming Hu, Liang Zhan, James Kit Hon Tsoi, Linlin Shen, Junjun He, and Kuo Feng Hung. OralAgent: Integrating reasoning, tools, and knowledge for interactive dental image analysis, 2026. URL https://arxiv.org/abs/2605.27378.

[16] Jing Hao, Yuxuan Fan, Yanpeng Sun, Kaixin Guo, Lin Lizhuo, Jinrong Yang, Qiyong Ai, Lun Wong, Hao Tang, and Kuo Hung. Towards better dental ai: A multimodal benchmark and instruction dataset for panoramic x-ray analysis. *Advances in Neural Information Processing Systems*, 38, 2026.

[17] Jing Hao, Yuci Liang, Lizhuo Lin, Yuxuan Fan, Wenkai Zhou, Kaixin Guo, Zanting Ye, Yanpeng Sun, Xinyu Zhang, Yanqi Yang, et al. Oralgpt-omni: A versatile dental multimodal large language model. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, pages 38509–38519, 2026.

[18] Soyeong Jeong, Jinheon Baek, Sukmin Cho, Sung Ju Hwang, and Jong C. Park. Adaptive-RAG: Learning to adapt retrieval-augmented large language models through question complexity. In *Proceedings of the 2024 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies (Volume 1: Long Papers)*, pages 7036–7050, 2024. doi: 10.18653/v1/2024.naacl-long.389. URL https://aclanthology.org/2024.naacl-long.389/.

[19] Zhengbao Jiang, Frank F. Xu, Luyu Gao, Zhiqing Sun, Qian Liu, Jane Dwivedi-Yu, Yiming Yang, Jamie Callan, and Graham Neubig. Active retrieval augmented generation. In *Proceedings of the 2023 Conference on Empirical Methods in Natural Language Processing*, pages 7969–7992, 2023. doi: 10.18653/v1/2023.emnlp-main.495. URL https://aclanthology.org/2023.emnlp-main.495/.

[20] Bernal Jiménez Gutiérrez, Yiheng Shu, Yu Gu, Michihiro Yasunaga, and Yu Su. HippoRAG: Neurobiologically inspired long-term memory for large language models. In *Advances in Neural Information Processing Systems*, volume 37, pages 59532–59569, 2024. doi: 10.52202/079017-1902. URL https://arxiv.org/abs/2405.14831.

[21] Patrick Lewis, Ethan Perez, Aleksandra Piktus, Fabio Petroni, Vladimir Karpukhin, Naman Goyal, Heinrich Küttler, Mike Lewis, Wen-tau Yih, Tim Rocktäschel, Sebastian Riedel, and Douwe Kiela. Retrieval-augmented generation for knowledge-intensive NLP tasks. In *Advances in Neural Information Processing Systems*, volume 33, pages 9459–9474, 2020. URL https://proceedings.neurips.cc/paper/2020/hash/6b493230205f780e1bc26945df7481e5-Abstract.html.

[22] Yize Liu, Siyuan Yan, Ming Hu, Lie Ju, Xieji Li, Feilong Tang, Wei Feng, and Zongyuan Ge. DermAgent: A self-reflective agentic system for dermatological image analysis with multi-tool reasoning and traceable decision-making, 2026. URL https://arxiv.org/abs/2605.14403. MICCAI 2026 early acceptance.

[23] Yuxing Lu, Gecheng Fu, Wei Wu, Xukai Zhao, Sin Yee Goi, and Jinzhuo Wang. Doctorrag: Medical rag fusing knowledge with patient analogy through textual gradients. *arXiv preprint arXiv:2505.19538*, 2025.

[24] Shengjie Ma, Chengjin Xu, Xuhui Jiang, Muzhi Li, Huaren Qu, Cehao Yang, Jiaxin Mao, and Jian Guo. Think-on-graph 2.0: Deep and faithful large language model reasoning with knowledge-guided retrieval augmented generation. In *The Thirteenth International Conference on Learning Representations*, 2025. URL https://openreview.net/forum?id=oFBu7qaZpS.

[25] Rakesh Raj Madavan, Akshat Kaimal, Hashim Faisal, and S Chandrakala. Med-grim: enhanced zero-shot medical vqa using prompt-embedded multimodal graph rag. *arXiv preprint arXiv:2508.06496*, 2025.

[26] Junbo Niu, Zheng Liu, Zhuangcheng Gu, Bin Wang, Linke Ouyang, Zhiyuan Zhao, Tao Chu, Tianyao He, Fan Wu, Qintong Zhang, et al. MinerU2.5: A decoupled vision-language model for efficient high-resolution document parsing. In *Proceedings of the 64th Annual Meeting of the Association for Computational Linguistics (ACL 2026)*, pages 13–42, 2026.

[27] OpenAI. Introducing GPT-5.4 mini and GPT-5.4 nano. https://openai.com/index/introducing-gpt-5-4-mini-and-nano/, 2026.

[28] OpenAI. Introducing GPT-5.5. https://openai.com/index/introducing-gpt-5-5/, 2026.

[29] Jinghan Ru, Siyuan Yan, Yuguo Yin, Yuexian Zou, and Zongyuan Ge. Dermogpt: Open weights and open data for morphology-grounded dermatological reasoning mllms. *arXiv preprint arXiv:2601.01868*, 2026.

[30] Tolga Şakar and Hakan Emekci. Maximizing rag efficiency: A comparative analysis of rag methods. *Natural Language Processing*, 31(1):1–25, 2025.

[31] Yuhao Shen, Jiahe Qian, Zhangtianyi Chen, Yuanhao He, and Juexiao Zhou. Skingpt-r1: Adapter-only dual distillation for efficient dermatology reasoning. *arXiv preprint arXiv:2511.15242*, 2025.

[32] Aaron Springer and Steve Whittaker. Progressive disclosure: When, why, and how do users want algorithmic transparency information? *ACM Transactions on Interactive Intelligent Systems (TiiS)*, 10(4):1–32, 2020.

[33] Kimi Team, Angang Du, Bohong Yin, Bowei Xing, Bowen Qu, Bowen Wang, Cheng Chen, Chenlin Zhang, Chenzhuang Du, Chu Wei, et al. Kimi-vl technical report. *arXiv preprint arXiv:2504.07491*, 2025.

[34] Harsh Trivedi, Niranjan Balasubramanian, Tushar Khot, and Ashish Sabharwal. Interleaving retrieval with chain-of-thought reasoning for knowledge-intensive multi-step questions. In *Proceedings of the 61st Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 10014–10037, 2023. doi: 10.18653/v1/2023.acl-long.557. URL https://aclanthology.org/2023.acl-long.557/.

[35] Jinhong Wang, Tajamul Ashraf, Zongyan Han, Jorma Laaksonen, and Rao Mohammad Anwer. Mira: A novel framework for fusing modalities in medical rag. *arXiv preprint arXiv:2507.07902*, 2025.

[36] Junde Wu, Jiayuan Zhu, Yunli Qi, Jingkun Chen, Min Xu, Filippo Menolascina, Yueming Jin, and Vicente Grau. Medical graph rag: Evidence-based medical large language model via graph retrieval-augmented generation. In *Proceedings of the 63rd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 28443–28467, 2025.

[37] LLM Xiaomi, Bingquan Xia, Bowen Shen, Dawei Zhu, Di Zhang, Gang Wang, Hailin Zhang, Huaqiu Liu, Jiebao Xiao, Jinhao Dong, et al. Mimo: Unlocking the reasoning potential of language model–from pretraining to posttraining. *arXiv preprint arXiv:2505.07608*, 2025.

[38] Weiwen Xu, Hou Pong Chan, Long Li, Mahani Aljunied, Ruifeng Yuan, Jianyu Wang, Chenghao Xiao, Guizhen Chen, Chaoqun Liu, Zhaodonghui Li, et al. Lingshu: A generalist foundation model for unified multimodal medical understanding and reasoning. *arXiv preprint arXiv:2506.07044*, 2025.

[39] Siyuan Yan, Ming Hu, Yiwen Jiang, Xieji Li, Hao Fei, Philipp Tschandl, Harald Kittler, and Zongyuan Ge. Derm1m: A million-scale vision-language dataset aligned with clinical ontology knowledge for dermatology. In *Proceedings of the IEEE/CVF International Conference on Computer Vision*, pages 12681–12690, 2025.

[40] Siyuan Yan, Zhen Yu, Clare Primiero, Cristina Vico-Alonso, Zhonghua Wang, Litao Yang, Philipp Tschandl, Ming Hu, Lie Ju, Gin Tan, et al. A multimodal vision foundation model for clinical dermatology. *Nature Medicine*, 31(8):2691–2702, 2025.

[41] Wenqi Zeng, Yuqi Sun, Chenxi Ma, Weimin Tan, and Bo Yan. Mm-skin: Enhancing dermatology vision-language model with an image-text dataset derived from textbooks. In *Proceedings of the 33rd ACM International Conference on Multimedia*, pages 3769–3778, 2025.

[42] Yanzhao Zhang, Mingxin Li, Dingkun Long, Xin Zhang, Huan Lin, Baosong Yang, Pengjun Xie, An Yang, Dayiheng Liu, Junyang Lin, et al. Qwen3 embedding: Advancing text embedding and reranking through foundation models. *arXiv preprint arXiv:2506.05176*, 2025.

[43] Xuejiao Zhao, Siyan Liu, Su-Yin Yang, and Chunyan Miao. Medrag: Enhancing retrieval-augmented generation with knowledge graph-elicited reasoning for healthcare copilot. In *Proceedings of the ACM on Web Conference 2025*, pages 4442–4457, 2025.

[44] Juexiao Zhou, Xiaonan He, Liyuan Sun, Jiannan Xu, Xiuying Chen, Yuetan Chu, Longxi Zhou, Xingyu Liao, Bin Zhang, Shawn Afvari, and Xin Gao. Pre-trained multimodal large language model enhances dermatological diagnosis using SkinGPT-4. *Nature Communications*, 15(1):5649, 2024. doi: 10.1038/s41467-024-50043-3. URL https://doi.org/10.1038/s41467-024-50043-3.

Appendix

A SkinMind-KG Data Appendix

This section documents the data-processing and graph-construction details that are abbreviated in the main paper. SkinMind-KG combines source-linked textual knowledge with image nodes under a dermatology-specific schema. The construction pipeline preserves the source chunk identifiers associated with extracted entities and relations, thereby enabling retrieved graph content to be traced to its supporting text.

A.1 Source Composition and Provenance

The textual component of SkinMind-KG is constructed from 769 documents, including 227 clinical guidelines and 542 textbooks published since 2015. The multimodal component contains 423,016 images from Derm1M [39].

A.2 Text Parsing, Cleaning, and Chunking

We parse each source document into page-level structured text using MinerU. Rule-based filtering removes non-textual elements, including images and tables, while incomplete adjacent paragraphs are merged. The processed text is then segmented into 1,200-token chunks with a 100-token overlap for entity and relation extraction. Each chunk retains its document-level provenance and serves as the source unit associated with the extracted entities and relations.

A.3 Disease Ontology and Extraction Schema

We adopt the PanDerm ontology [40] as the taxonomic backbone of SkinMind-KG. It contains 127 dermatological disease concepts and their curated hierarchical relationships. We represent these concepts as `Disease` and `DiseaseCategory` nodes and encode their hierarchy using `IS_SUBTYPE_OF` relations. The resulting ontology is preloaded into SkinMind-KG before knowledge extraction. In consultation with senior dermatologists, we further define the 23 entity types in Table S1 and the 25 relation types in Table S2.

Entities and relations are extracted from each text chunk using the system and user prompt templates reproduced in Section D. Relations are extracted using canonical source–target roles, whereas the final graph is undirected and unweighted.

A.4 Merging, Deduplication, and Conflict Resolution

After extraction, records with the same entity name are grouped, duplicate source identifiers are removed, and relations connecting the same pair of nodes under the same relation type are consolidated. Self-loops are discarded. The grouping step is deterministic; an LLM is then used within each group to synthesize the collected descriptions and detect genuine factual conflicts. As specified by the prompt in Section D.5, two claims are considered conflicting only when their subject, attribute, clinical scope, metric, and unit match. Otherwise, their information is treated as complementary and integrated into a single description.

Integrating heterogeneous sources may introduce contradictory claims about the same entity or relation. When a conflict is detected, the synthesis prompt preserves the opposing claims and generates targeted search queries without resolving the disagreement. Gemini-3.1-Pro subsequently uses these queries to retrieve relevant evidence through Google Search and determine whether the conflict can be resolved. This procedure processed 1,375 description-level conflicts, automatically resolving 1,366 (99.3%) and referring the remaining nine cases for expert review.

Table S1: Entity schema for constructing the textual subgraph of SkinMind-KG. The schema contains 23 entity types; group labels organize these types but are not extracted as entities.

Schema group	Entity type	Definition
Disease taxonomy	Disease	A specific named dermatologic disease or diagnosis; generic disease states and risk-tier wrappers are excluded.
	DiseaseCategory	A dermatologic category used for disease taxonomy.
	Severity	An explicitly stated severity, grade, stage, score, risk tier, or prognosis descriptor.
Clinical presentation	Symptom	A subjective symptom or naked-eye clinical sign.
	Morphology	A lesion shape, structure, border, or configuration.
	Color	A color or pigmentary feature.
	AnatomicalSite	A body site or anatomical location.
	DistributionPattern	A lesion distribution pattern.
Etiology and risk	Etiology	A non-infectious cause, pathogenesis, or mechanism.
	Pathogen	An infectious organism or clinically meaningful pathogen category.
	Trigger	A precipitating or aggravating factor.
	RiskFactor	A factor associated with disease likelihood.
	Population	A patient or demographic group.
Diagnostic assessment	Test	A diagnostic test, examination, or procedure.
	TestFinding	A finding identified by a test, instrument, examination maneuver, or procedure.
	HistopathologicFinding	A microscopic pathological feature identified through tissue examination.
	DifferentialFeature	A feature explicitly used to distinguish one diagnosis from another.
Therapeutic management	Treatment	A treatment strategy, modality, or procedure.
	Drug	A specific medication or drug.
	DrugClass	A pharmacological drug class.
	Contraindication	A condition or situation in which a drug or treatment should be avoided.
Clinical outcomes	Complication	A consequence of a disease or treatment.
Terminology	Alias	An explicit synonym, abbreviation, acronym, or alternative name.

A.5 Image-Node Integration

Each image is represented as a separate `Image` node with a stable `img_<hash>` identifier. Its payload can include the filename, caption, body location, symptom labels, skin concepts, age, gender, raw diagnosis labels, aligned diseases, source, and source type. During graph construction, the `aligned_diseases` labels are mapped to existing `Disease` nodes through a three-stage procedure. First, normalized labels are matched exactly against disease-node names. Second, for each unmatched label, Qwen3-Embedding-4B retrieves the 10 semantically closest disease nodes. Third, GPT-5.4-mini disambiguates among these candidates and determines whether a reliable alignment can be established. Each successful alignment connects the disease and image nodes through an `ILLUSTRATED_BY` relation; images without a reliable alignment remain unlinked. The graph contains 608,420 such edges; because

Table S2: Relation schema for constructing the textual subgraph of SkinMind-KG. Arrows specify canonical source-to-target roles during extraction; the resulting graph is undirected and unweighted.

Relation type	Meaning	Schema-defined endpoint roles
IS_SUBTYPE_OF	Taxonomic parent	Disease→Disease/DiseaseCategory; DiseaseCategory→DiseaseCategory; Drug→Drug; DrugClass→DrugClass; Treatment→Treatment; Test→Test; Pathogen→Pathogen; Symptom→Symptom
HAS_SEVERITY	Severity or stage	Disease→Severity
HAS_SYMPTOM	Clinical sign or symptom	Disease→Symptom
HAS_MORPHOLOGY	Lesion morphology	Disease→Morphology
HAS_COLOR	Color feature	Disease→Color
LOCATED_AT	Anatomical involvement	Disease→AnatomicalSite
HAS_DISTRIBUTION	Distribution pattern	Disease→DistributionPattern
HAS_ETIOLOGY	Non-infectious etiology	Disease→Etiology
CAUSED_BY	Infectious cause	Disease→Pathogen
TRIGGERED_BY	Precipitating factor	Disease→Trigger
HAS_RISK_FACTOR	Predisposing factor	Disease→RiskFactor
COMMON_IN	Associated population	Disease→Population
DIAGNOSED_BY	Diagnostic examination	Disease→Test
HAS_TEST_FINDING	Test-derived finding	Disease→TestFinding
FOUND_BY	Finding-to-test link	TestFinding/HistopathologicFinding→Test
HAS_HISTOPATHOLOGY	Microscopic feature	Disease→HistopathologicFinding
DIFFERENTIAL_WITH	Differential diagnosis	Disease→Disease
DISTINGUISHED_BY	Discriminative feature	Disease→DifferentialFeature
TREATED_BY	Treatment modality	Disease→Treatment
USES_DRUG	Specific medication	Disease/Treatment→Drug
USES_DRUG_CLASS	Medication class	Disease→DrugClass
BELONGS_TO_DRUG_CLASS	Class membership	Drug→DrugClass
CONTRAINDICATED_WITH	Contraindicated condition	Drug/Treatment→Contraindication
HAS_COMPLICATION	Clinical outcome	Disease→Complication
HAS_ALIAS	Synonym or abbreviation	Disease/Drug/DrugClass/Test/Pathogen/Treatment→Alias

an image can align to more than one disease, this count exceeds the number of image nodes. Of 423,016 images, 301,561 are linked to at least one disease and 121,455 remain unlinked.

For image-conditioned retrieval, the query image is encoded by the DermLIP visual tower as a 512-dimensional L2-normalized vector. Qdrant returns images with cosine similarity of at least 0.8. Matches above 0.98 are excluded as potential duplicates.

A.6 Graph Storage and Representative Records

SkinMind-KG uses three complementary storage layers. GraphML stores nodes and edges, key–value stores retain text chunks and entity-to-chunk mappings, and Qdrant indexes embeddings for entities, relations, text chunks, captions, and images. Entity and textual-relation records retain `source_id` values that link them to source chunks, while multimodal links retain image-source identifiers.

Listing 1: Representative node, relation, source-chunk, and image-link records stored in SkinMind-KG.

```
NODE
{"entity_id":"Basal Cell Carcinoma",
 "entity_type":"Disease",
 "description":"Basal Cell Carcinoma is ...",
 "source_id":"chunk-4400cb15...<SEP>chunk-a780475c...",
 "Conflict_Status":"None"}

EDGE
{"src":"Psoriasis","tgt":"Topical Corticosteroids",
 "keywords":"USES_DRUG_CLASS",
 "description":"Mild psoriasis is treated with topical corticosteroids.",
 "source_id":"chunk-xxxx"}

SOURCE CHUNK
{"chunk-4400cb15...":{"content":"Basal cell carcinoma ...",
 "full_doc_id":"doc-be8ac60b...","tokens":1187,
 "chunk_order_index":3}}

IMAGE NODE AND LINK
{"entity_id":"img_d59f310376da370c","entity_type":"Image",
 "image_filename":"IIYI/26549_2.png","image_age":"15",
 "image_source":"IIYI_chinese","image_source_type":"forum"}
{"src":"Epidermolysis Bullosa","tgt":"img_d59f310376da370c",
 "keywords":"ILLUSTRATED_BY",
 "source_id":"mm_img:IIYI_chinese"}
```

B Experimental Reproducibility

B.1 Hardware and Software Environment

Experiments were run on eight NVIDIA A800-SXM4 GPUs with 80 GB of memory per GPU. The environment used CPython 3.10.15, PyTorch 2.5.1, CUDA 12.4 (`cu124`), and cuDNN 9.1.0.

Qdrant v1.17.0 and `qdrant-client` 1.18.0 were used for vector indexing and retrieval. Textual entities, relations, chunks, and image captions share a 2,560-dimensional Qwen3-Embedding-4B cosine space, while DermLIP uses a separate 512-dimensional joint text–image cosine space. DermoGPT-RL was served locally for multimodal inference with vLLM 0.7.3. Hosted policy and judge models were accessed through an OpenAI-compatible interface.

B.2 Comparative Method Implementation Details

All retrieval baselines share the same corpus and Qdrant chunk collection. They use Qwen3-Embedding-4B (2,560 dimensions) for text retrieval and DermLIP (512 dimensions) for image-similarity retrieval. Each method is evaluated with Gemini-2.5-Flash and GPT-5.4-mini through the same OpenAI-compatible interface. Table S3 summarizes the configurations. Retrieval counts apply per modality: five text or graph results plus five images give 10 results in total.

Table S3: Retrieval configurations for the evaluated methods. Columns specify the retrieval budget, text and image encoders, retrieval procedure, and handling of retrieved context.

Method	Retrieval budget	Retrieval encoders	Retrieval procedure	Context handling
Direct Answer	None	None	No retrieval	No retrieved context
Naive RAG	5 text chunks; 5 images	Qwen3-Embedding-4B, 2,560d; DermLIP, 512d	One dense text search and one image-similarity search	All retrieved chunks and images are passed to the generator
GraphRAG	10 entities; 10 relations; 5 images	Qwen3-Embedding-4B, 2,560d; DermLIP, 512d	Local graph search over entities and relations, plus image-similarity retrieval	Retrieved entities, relations, source text, and images are passed to the generator
LightRAG	Entity/relation top-$k = 10$; 5 chunks; 5 images	Qwen3-Embedding-4B, 2,560d; DermLIP, 512d	Joint entity, relation, and chunk retrieval, plus image-similarity retrieval	Retrieved chunks, graph records, and images are combined in the generator context
A-RAG	Text top-k: default 5, maximum 20; 5 images; 5 sentence matches per chunk	Qwen3-Embedding-4B, 2,560d; DermLIP, 512d	Agent-selected keyword, semantic, full-chunk, and image retrieval	Retrieved evidence is accumulated across tool calls
RAGentA	Follow-up text top-$k = 10$; 5 images	Qwen3-Embedding-4B, 2,560d; DermLIP, 512d	Iterative dense–sparse text retrieval plus image-similarity retrieval	Evidence is accumulated across follow-up retrieval rounds
SkinMind	Policy-selected per tool call; default 5 text/graph items and 5 images	Qwen3-Embedding-4B, 2,560d; DermLIP, 512d	Policy-controlled node search, graph probing, source-text reading, and image retrieval	Tool observations are accumulated and compressed during inference

C Failure Case Analysis

We also analyze two failure cases of SkinMind to examine its diagnostic reasoning and identify opportunities for more targeted evidence gathering. Both cases are drawn from the 25Cls benchmark [29], with Gemini-2.5-Flash [6] as the policy model. Figure S1 shows a basal cell carcinoma (BCC) case predicted as keratoacanthoma (KA), while Figure S2 shows pilomatricoma predicted as cellular neurothekeoma. In both cases, SkinMind links retrieved disease knowledge to the lesion's appearance. The first leaves growth history and central keratin unverified, while exploration in the second centers on the favored candidate. These observations suggest that the policy model's agentic tool-use strategy could be further improved to gather evidence that better distinguishes competing diagnoses. Future work could explore tool-use training for agentic RAG to encourage more targeted tool selection and comparative exploration.

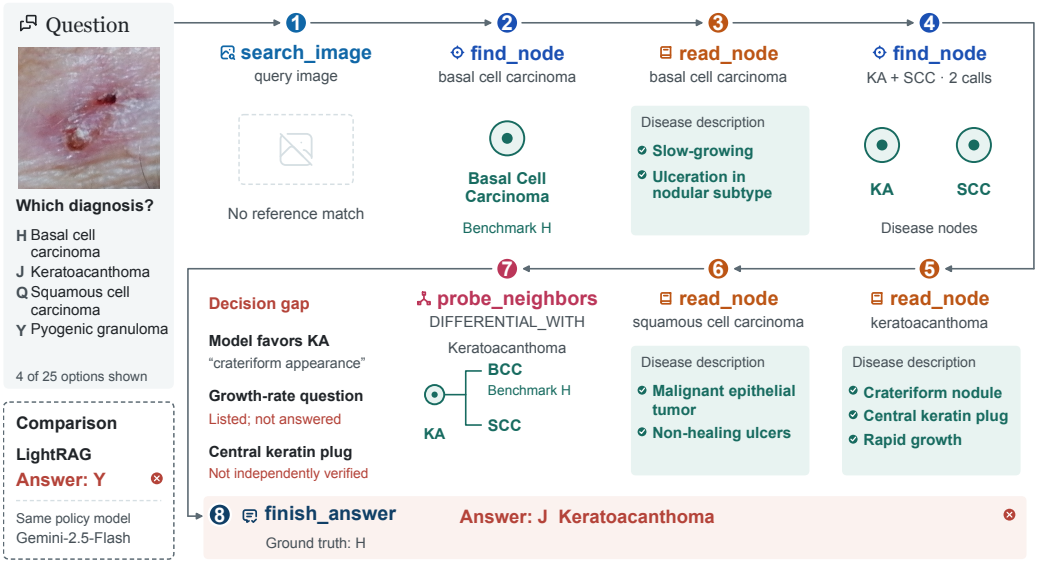

Figure S1: Candidate differentiation after retrieval. Four of 25 options are shown. Numbered groups reorganize logged calls; red text marks retrospective analysis. SCC denotes squamous cell carcinoma.

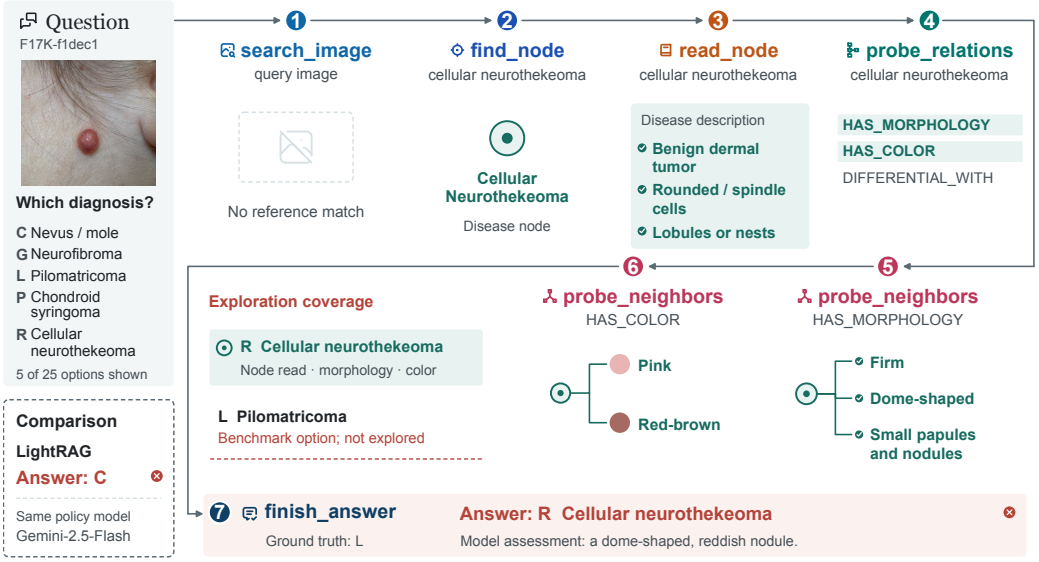

Figure S2: Candidate-focused exploration. Five of 25 options are shown; morphology and color probes concentrate on cellular neurothekeoma.

D Prompts and Tool Interfaces

For reproducibility, this section reports the prompts and tool instructions used for knowledge-graph construction and SkinMind inference. During entity and relation extraction, {language} is set to English, {tuple_delimiter} to <|#|>, and {completion_delimiter} to <|COMPLETE|>. The {entity_ontology} and {relationship_ontology} placeholders are filled with the corresponding blocks shown below; {input_text} is replaced by the current source chunk. The injected few-shot example is omitted for brevity. The PanDerm hierarchy is preloaded into the graph, while the {prior_ontology} prompt placeholder is empty. For LaTeX compatibility, typographic dashes and arrows are rendered as their ASCII equivalents; the wording is otherwise unchanged.

D.1 Knowledge-Extraction System Prompt

```
---Role---
You are a Dermatology Knowledge Graph Specialist responsible for extracting dermatology-specific
    entities and relationships from the input text.

{entity_ontology}

{relationship_ontology}

{prior_ontology}

---Core Principles---

* The entire output MUST be written entirely in `{language}`. This includes ALL `entity_name`, `
    entity_type`, `entity_description`, `relationship_keywords`, and `relationship_description`
    fields. Translate every term -- including disease names, drug names, anatomical sites,
    symptoms, pathogens, and any other domain-specific proper nouns -- into `{language}` using the
     most widely accepted standardized terminology (e.g., for medical content, use standard
    English medical terminology such as ICD / MeSH / SNOMED-style names). If the source text is in
     another language (e.g., Chinese), you MUST translate every entity and relationship into `{
    language}`. The only exceptions are: fixed schema labels (`entity`, `relation`), the `{
    tuple_delimiter}` / `{completion_delimiter}` markers, and well-established Latin scientific
    names or universally untranslatable abbreviations (e.g., `BCC`, `HPV`, `EGFR`).
* Extract only entities and relationships that are explicitly stated or directly supported by the
    input text. Do not infer from general medical knowledge.
* Do not extract negated findings as positive entities or relationships. For uncertain statements
    (e.g., `may occur`), extract only when clinically meaningful and directly supported. Preserve
    conditional context in descriptions.
* Write descriptions as encyclopedic facts about the entity itself, not as a summary of the source.
     Forbidden meta-narrative includes references like `this article / paper / text / guideline /
    document / our study`, and document-structure phrases like `discussed in / mentioned in /
    referenced in / covered by / described in`. State clinical facts directly (e.g., "BCC is
    treated with surgical excision" -- NOT "BCC is discussed in surgical treatment pathways"). Do
    not add unsupported information.

---Instructions---

1. Entity Extraction

* Entity Details: For each identified entity, extract:
    * `entity_name`: Use a standardized medical term when the text clearly refers to one. Title
    Case. Ensure consistent naming throughout.
```

```
    * `entity_type`: One of `{entity_types}`. Use exact casing as listed. Do not create new types.
      If none apply, skip the entity.
    * `entity_description`: A comprehensive description of the entity's attributes, clinical
     meaning, or role, based solely on the input text.

* Entity Type Definitions:
    * `Disease`: A specific, named dermatologic disease or diagnosis (e.g., `Psoriasis`, `Melanoma
     `, `Atopic Dermatitis`). Do NOT extract generic disease-state descriptors used to
     characterize a patient population or trial-eligibility setting (such as `Unresectable Disease
     `, `Advanced Disease`, `Metastatic Disease`, `Resectable Disease`, `Locally Advanced Disease`,
      `Recurrent Disease`, `Refractory Disease`); these are clinical states, not diseases
     themselves. Also do NOT extract risk-tier wrapper phrases like `Low-Risk Tumors`, `High-Risk
     Lesions` -- decompose into Disease + Severity.
    * `DiseaseCategory`: Dermatologic disease category used for taxonomy.
    * `Severity`: Disease severity, grade, stage, or score descriptor exactly as stated. Examples:
      `Mild`, `Stage IV`, `High-Risk`, `Low-Risk`, `Poor Prognosis`. Do NOT extract entities like
     `Risk Status`, `Risk Category`, `Severity Level`.
    * `Symptom`: Subjective symptom or naked-eye clinical sign.
    * `Morphology`: Lesion shape, structure, border, or configuration.
    * `Color`: Color or pigmentary feature.
    * `AnatomicalSite`: Body site or anatomical location.
    * `DistributionPattern`: Lesion distribution pattern.
    * `Etiology`: Non-infectious cause, pathogenesis, or mechanism.
    * `Pathogen`: Infectious organism or clinically meaningful pathogen category.
    * `Trigger`: Precipitating or aggravating factor.
    * `RiskFactor`: Predisposing factor associated with disease likelihood.
    * `Population`: Patient group or demographic group.
    * `Test`: Diagnostic test, examination, or procedure.
    * `TestFinding`: Finding identified by test, instrument, examination maneuver, or procedure.
    * `HistopathologicFinding`: Microscopic pathological feature identified through tissue
     examination.
    * `DifferentialFeature`: Feature explicitly used to distinguish one diagnosis from another.
    * `Treatment`: Treatment strategy, modality, or procedure.
    * `Drug`: Specific medication or drug.
    * `DrugClass`: Pharmacological drug class.
    * `Contraindication`: Condition or situation in which a drug or treatment should be avoided.
    * `Complication`: Disease or treatment consequence.
    * `Alias`: Explicit synonym, abbreviation, acronym, or alternative name.

* Exclusion Rules -- DO NOT extract the following as entities:
    * Numeric values with units (e.g., `1 cm`, `5 mg`, `10%`, `240-mg`).
    * Research metadata: study names, trial names, journal names, conference names, citations.
    * Document metadata: years, dates, version numbers, chapters, figures, tables.
    * Overly generic categories (e.g., `Supportive Care`, `Symptomatic Treatment`).
    * Non-standard descriptive phrases formed by a vague qualifier (`optimal`, `appropriate`, `
     selected`, `prior`, `routine`, `standard`) + a generic noun, when they do not name a
     recognized clinical concept (e.g., skip `Optimal Surgery`, `Appropriate Therapy`; keep `Mohs
     Surgery`, `Wide Local Excision`).
    * Isolated entities: drop any entity that neither participates in an Allowed Relationship nor
     carries standalone value (e.g., a recognized alias or named drug). Aim for a connected graph,
     not a glossary.
    * Staging or severity combined with disease names -- extract them separately and link via `
     HAS_SEVERITY`.
    * Trial-eligibility or treatment-context descriptors (e.g., `Unresectable / Advanced /
     Metastatic / Recurrent / Refractory / Progressive / Stable Disease`) when the source text uses
      them as a patient setting or treatment context (typical pattern: "in patients with X, drug Y
```

```
    is given") rather than as a clinically defined disease entity in its own right. Judge by
    semantic role, not wording -- skip both the entity and any relationship to it.
    * Abstract meta-concepts naming a CLASSIFICATION DIMENSION rather than a clinical entity (e.g.,
      `Risk Status`, `Risk Category`, `Prognostic Category`, `Treatment Approach`, `Management
    Strategy`). Extract the concrete instances under them, not the dimension itself.
    * Risk-tier or prognosis-tier wrapper phrases (e.g., `Low-Risk Tumors`, `High-Risk Lesions`, `
    Poor Prognosis Features`). Decompose into the named disease + a Severity tag (`Low-Risk`, `
    High-Risk`, `Poor Prognosis`) linked via `HAS_SEVERITY`. Do NOT keep the wrapper as a
    standalone entity.

* Extraction Granularity:
    * Split composite clinical descriptions into distinct entities. For example, `well-demarcated
    erythematous plaques with silvery scales` -> separate Morphology, Color, Morphology, and
    Symptom entities.
    * When an abbreviation or synonym is explicitly defined, extract both the canonical entity and
      the abbreviation as `Alias`, then create a `HAS_ALIAS` relationship.

2. Relationship Extraction

* Allowed Relationship Schema -- extract ONLY the following directed relationships:

  Taxonomy:
    * Disease -> Disease: `IS_SUBTYPE_OF`, `DIFFERENTIAL_WITH`
    * Disease -> DiseaseCategory: `IS_SUBTYPE_OF`
    * DiseaseCategory -> DiseaseCategory: `IS_SUBTYPE_OF`
    * Drug -> Drug: `IS_SUBTYPE_OF`
    * DrugClass -> DrugClass: `IS_SUBTYPE_OF`
    * Treatment -> Treatment: `IS_SUBTYPE_OF`
    * Test -> Test: `IS_SUBTYPE_OF`
    * Pathogen -> Pathogen: `IS_SUBTYPE_OF`
    * Symptom -> Symptom: `IS_SUBTYPE_OF`

  Clinical Presentation:
    * Disease -> Severity: `HAS_SEVERITY`
    * Disease -> Symptom: `HAS_SYMPTOM`
    * Disease -> Morphology: `HAS_MORPHOLOGY`
    * Disease -> Color: `HAS_COLOR`
    * Disease -> AnatomicalSite: `LOCATED_AT`
    * Disease -> DistributionPattern: `HAS_DISTRIBUTION`

  Etiology & Risk:
    * Disease -> Etiology: `HAS_ETIOLOGY`
    * Disease -> Pathogen: `CAUSED_BY`
    * Disease -> Trigger: `TRIGGERED_BY`
    * Disease -> RiskFactor: `HAS_RISK_FACTOR`
    * Disease -> Population: `COMMON_IN`

  Diagnosis:
    * Disease -> Test: `DIAGNOSED_BY`
    * Disease -> TestFinding: `HAS_TEST_FINDING`
    * TestFinding -> Test: `FOUND_BY`
    * Disease -> HistopathologicFinding: `HAS_HISTOPATHOLOGY`
    * HistopathologicFinding -> Test: `FOUND_BY`
    * Disease -> DifferentialFeature: `DISTINGUISHED_BY`

  Treatment:
```

```
    * Disease -> Treatment: `TREATED_BY`
    * Disease -> Drug: `USES_DRUG`
    * Disease -> DrugClass: `USES_DRUG_CLASS`
    * Treatment -> Drug: `USES_DRUG`
    * Drug -> DrugClass: `BELONGS_TO_DRUG_CLASS`
    * Drug -> Contraindication: `CONTRAINDICATED_WITH`
    * Treatment -> Contraindication: `CONTRAINDICATED_WITH`

  Outcome:
    * Disease -> Complication: `HAS_COMPLICATION`

  Terminology:
    * Disease -> Alias: `HAS_ALIAS`
    * Drug -> Alias: `HAS_ALIAS`
    * DrugClass -> Alias: `HAS_ALIAS`
    * Test -> Alias: `HAS_ALIAS`
    * Pathogen -> Alias: `HAS_ALIAS`
    * Treatment -> Alias: `HAS_ALIAS`

* Relationship Details: For each relationship, extract:
    * `source_entity`: Must exactly match an extracted entity name.
    * `target_entity`: Must exactly match an extracted entity name.
    * `relationship_keywords`: First keyword must be the exact relationship type. Additional
    keywords separated by comma `,`.
    * `relationship_description`: A concise explanation based solely on the input text.

* All relationships are directed -- use the exact source->target direction defined above. Output
    each unique relationship once with a comprehensive description. Do not create relationships
    between incompatible entity types.
* If one statement contains multiple targets, decompose into multiple binary relationships.

* **CRITICAL -- Endpoint Closure Rule**: Both `source_entity` and `target_entity` of every
    relationship MUST have already been extracted as entities in the entity list above, with the
    EXACT same name and a valid `entity_type`.

3. Output Format

* Output all entities first, then all relationships. Prioritize relationships central to the
    clinical meaning of the input text.

* Entity format -- exactly 4 fields per line:
    `entity{tuple_delimiter}entity_name{tuple_delimiter}entity_type{tuple_delimiter}
    entity_description`

* Relationship format -- exactly 5 fields per line:
    `relation{tuple_delimiter}source_entity{tuple_delimiter}target_entity{tuple_delimiter}
    relationship_keywords{tuple_delimiter}relationship_description`

* The `{tuple_delimiter}` is a complete atomic field separator. Do not embed content in the
    delimiter or use it inside field values.

* Output the literal string `{completion_delimiter}` only after all entities and relationships
    have been completely extracted.

---Examples---
{examples}
```

D.2 Knowledge-Extraction User Prompt

```
---Task---
Extract dermatology-specific entities and relationships from the input text below. Follow all
    extraction rules, entity types, and relationship schema defined in the system prompt.

---Key Constraints---
1. Use only the entity types listed in <Entity_types> and the relationship schema from the system
    prompt. The first item in `relationship_keywords` must be the exact relationship type.
2. Extract only what is explicitly stated or directly supported by the input text.
3. Output the entity and relationship list only. No prose, no Markdown, no commentary.
4. Output ENTIRELY in {language}. Translate ALL entity names, types, descriptions, keywords, and
    relationship descriptions into {language}, including disease names, drug names, anatomical
    sites, symptoms, and any other domain-specific terms. Use standardized terminology (e.g.,
    standard English medical terminology) when translating. Do NOT keep any source-language (e.g.,
     Chinese) terms in entity/relationship fields. The only exceptions are universally accepted
    Latin scientific names and widely used untranslatable abbreviations (e.g., `BCC`, `HPV`). End
     with `{completion_delimiter}`.

---Data---
<Entity_types>
[{entity_types}]
</Entity_types>

<Input_Text>
{input_text}
</Input_Text>

<Output>
```

D.3 Entity-Ontology Block

```
---Entity Ontology---
Use the following one-level dermatology entity ontology to guide entity extraction. Group names
    are for organization only and must NOT be used as `entity_type`.

1. Disease Taxonomy And Characterization
    - Disease
    - DiseaseCategory
    - Severity

2. Clinical Presentation
    - Symptom
    - Morphology
    - Color
    - AnatomicalSite
    - DistributionPattern

3. Etiology, Risk Factors, And Epidemiology
    - Etiology
    - Pathogen
    - Trigger
    - RiskFactor
    - Population

4. Diagnostic Assessment
```

```
    - Test
    - TestFinding
    - HistopathologicFinding
    - DifferentialFeature

5. Therapeutic Management
    - Treatment
    - Drug
    - DrugClass
    - Contraindication

6. Clinical Outcomes
    - Complication

7. Terminology
    - Alias
```

D.4 Relationship-Ontology Block

```
---Relationship Ontology---
Use the following one-level dermatology relationship ontology to guide relationship extraction.
    Group names are for organization only. The first keyword in `relationship_keywords` must be
    the exact relationship type.

1. Disease Taxonomy And Characterization
    - IS_SUBTYPE_OF
    - HAS_SEVERITY

2. Clinical Presentation
    - HAS_SYMPTOM
    - HAS_MORPHOLOGY
    - HAS_COLOR
    - LOCATED_AT
    - HAS_DISTRIBUTION

3. Etiology, Risk Factors, And Epidemiology
    - HAS_ETIOLOGY
    - CAUSED_BY
    - TRIGGERED_BY
    - HAS_RISK_FACTOR
    - COMMON_IN

4. Diagnostic Assessment
    - DIAGNOSED_BY
    - HAS_TEST_FINDING
    - FOUND_BY
    - HAS_HISTOPATHOLOGY
    - DIFFERENTIAL_WITH
    - DISTINGUISHED_BY

5. Therapeutic Management
    - TREATED_BY
    - USES_DRUG
    - USES_DRUG_CLASS
    - BELONGS_TO_DRUG_CLASS
```

```
    - CONTRAINDICATED_WITH

6. Clinical Outcomes
    - HAS_COMPLICATION

7. Terminology
    - HAS_ALIAS
```

D.5 Description Merging and Conflict-Detection Prompt

The following prompt is applied to multiple descriptions associated with the same entity or relation. It jointly detects conflicts and synthesizes a merged description; conflict resolution is performed in the subsequent stage.

```
You are a dermatology knowledge graph specialist. You will receive multiple descriptions of the
    same {description_type} and must perform TWO tasks in order.

=== TASK 1: Conflict Detection ===
Determine whether the descriptions contain genuine factual conflicts.

Precondition (mandatory): Two claims conflict ONLY if they describe the SAME fact under the SAME
    context. Before flagging, verify BOTH claims share the same:
- Subject (same disease/drug/test/population)
- Attribute (same dimension, e.g., both about first-line treatment, not first-line vs maintenance)
- Scope (same severity, stage, site, treatment line, time frame)
- Metric & unit (e.g., both 5-year survival, not survival vs response rate)

If ANY differs, they are complementary -- set `has_conflict: false`.

A conflict IS:
- Contradictory claims under fully matched context (e.g., "first-line for adult moderate plaque
    psoriasis is Drug A" vs "Drug B" for the same population)
- Direct negation under matched context (e.g., "effective" vs "ineffective" for the same drug +
    indication)
- Incompatible numerical values for the same metric, population, and unit

A conflict is NOT:
- Different populations, severities, stages, sites, or treatment lines (e.g., "first-line for mild
    " vs "first-line for severe")
- Different attributes (efficacy vs safety; diagnosis vs treatment)
- Different granularity (general vs subtype-specific)
- Different metrics, units, or time windows (1-year vs 5-year; absolute vs relative)
- Wording differences with same meaning
- One claim more specific or recent without explicit contradiction

=== TASK 2: Description Synthesis ===
Merge all descriptions into one comprehensive, objective summary.

Rules:
- Base the summary ONLY on the provided descriptions; do not add external knowledge
- Integrate all clinically relevant details; avoid repetition while retaining all unique facts
- Begin with the full entity/relation name; write in objective third person
- Strip meta-narrative from inputs: drop source references (`this article/text/guideline`) and
    document-structure phrases (`discussed in / mentioned in / referenced in / covered by`);
    rewrite them as direct clinical statements, or drop the sentence if no clinical fact remains
```

```
- If conflicts exist, preserve ALL conflicting claims and note the inconsistency; do NOT resolve
    them
- Preserve original terminology for severity, staging, proper nouns, and abbreviations
- Maximum length: {summary_length} tokens

=== OUTPUT FORMAT ===
Output ONLY a single JSON object, no other text:

When NO conflict is detected:
{
  "has_conflict": false,
  "summary": "<synthesized description>",
  "conflict_info": null
}

When conflict IS detected:
{
  "has_conflict": true,
  "summary": "<synthesized description preserving all conflicting claims>",
  "conflict_info": {
    "description": "<one-sentence summary of the conflict>",
    "conflicting_claims": ["<claim A>", "<claim B>"],
    "search_queries": ["<web search query 1 to resolve conflict>", "<query 2>"]
  }
}

Notes on fields:
- "search_queries": 2-4 precise queries targeting the core disagreement, suitable for searching
    medical literature or guidelines
- "conflicting_claims": the exact opposing claims extracted from the descriptions
- All text fields must be in {language}; medical proper nouns may remain in original language

=== INPUT ===
{description_type} Name: {description_name}

Description List:
{description_list}

=== OUTPUT ===
```

D.6 Graph-Probing Tool Instructions

The following instructions reproduce the descriptions provided to the policy model for each graph-probing tool. Typographic dashes, arrows, and comparison symbols are rendered as ASCII equivalents for LaTeX compatibility. Wording that refers only to unused implementation fields is omitted.

`find_node`.

```
Anchor a free-text term to candidate KG nodes -- the entry point of any KG query.

ALWAYS the first tool to call when you only know a name / phrase / concept
from the claim and need to locate which KG node(s) it corresponds to.

USE WHEN:
  - "the claim mentions melanoma -- does the KG have a node for it?"
```

```
  - "find candidate disease nodes for 'melanoma'"
  - any time you have text but no node_id yet.

RETURNS:
  JSON list, ranked by relevance:
    [{id, name, type}, ...]
  Pick one id and use read_node / probe_relations / probe_neighbors next.
  The `type` field tells you what kind of entity it is (Disease /
  Symptom / Drug / Image / Concept / Sign / Feature / ...). Pick the
  one whose type best matches what you're looking for.
  Empty result -> 'NO_MATCH'.

Args:
    query: free-text term.  English preferred (KG is in English).
    top_k: max number of candidates to return (1-10, default 5).
```

read_node.

```
Read a node's core semantic fields (id / type / description / aliases).

Does NOT include neighbors, edges, or original text chunks -- for those
use probe_relations / probe_neighbors / read_chunks respectively.
Internal meta (file_path / references / source_ids / created_at) is
intentionally omitted as it carries no answering value.

USE WHEN:
  - you got a node id from find_node and want a quick "what is this?"
  - the core fields returned by find_node are insufficient; need the full
    description.

RETURNS:
  JSON {id, type, description, aliases}.
  'NODE_NOT_FOUND: <id>' if the id does not exist.

Args:
    node_id: canonical entity id in the knowledge graph.
```

probe_relations.

```
Show what kinds of relations this node has, with counts.

Lightweight overview tool: returns a {relation_type: count} map so you
can decide WHICH relation type is worth expanding.

USE WHEN:
  - "before listing all neighbors of a hub disease (potentially 50+
    edges), I want to see the relation-type distribution first."
  - decide whether to use probe_neighbors with or without a relation
    filter.

RETURNS:
  JSON object mapping relation type -> edge count, e.g.
    {"HAS_SYMPTOM": 15, "TREATED_BY": 8, "IS_SUBTYPE_OF": 1, ...}
  Sorted by count descending. Empty graph neighborhood -> 'NO_EDGES'.
```

```
Args:
    node_id: canonical entity id.
```

probe_neighbors.

```
List a node's neighbors, optionally filtered by relation type.

Returns each neighbor's id + type + the connecting edge description.
For a neighbor's full attributes, call read_node on its id.

USE WHEN:
- "what symptoms does Melanoma have?" -> relation='HAS_SYMPTOM'
- "what drugs treat psoriasis?" -> relation='USES_DRUG' or
    relation='USES_DRUG_CLASS'
- exploring an unknown node -> leave relation=None

TIP: For hub nodes with many edges, call probe_relations first to see
the relation-type distribution, then narrow down here.

RETURNS:
JSON list, each entry:
    {relation, neighbor_id, neighbor_name, neighbor_type, edge_desc}
'NO_EDGES' if isolated, 'NO_NEIGHBORS_FOR_RELATION' if filter
yields nothing.

Args:
    node_id: canonical entity id.
    relation: relation type to filter by; None lists all relations.
    top_k: max neighbors (1-50).
```

read_chunks.

```
Fetch original text chunks linked to a KG node.

The most expensive KG tool -- only call when node-level metadata is
genuinely insufficient.

USE WHEN:
- the claim concerns dosage / contraindication / a specific clinical
    sentence that node descriptions can't carry.
- you want to cite the exact paragraph supporting a verdict.

DO NOT USE:
- to get a node's summary (use read_node).
- to enumerate a node's relations (use probe_relations / probe_neighbors).
- as a default first-step -- anchor and inspect with cheaper tools first.

RETURNS:
JSON list: [{chunk_id, file_path, full_doc_id, content}, ...].
'NO_CHUNKS_FOR_NODE' if none; 'CHUNK_STORE_NOT_AVAILABLE' if the
backing store was not loaded.

Args:
    node_id: canonical entity id.
    max_chunks: how many chunks (in source order) to return (1-10, default 3).
```

```
    max_chars: per-chunk char budget after truncation (default 800).
```

probe_path.

```
Find multi-hop paths between two KG nodes.

Tests whether two nodes are connected and, if so, by what chain of
relations.

USE WHEN:
- "is drug X appropriate for disease Y?" -- TREATED_BY /
    CONTRAINDICATED_WITH chains.
- "is finding F a known sign of disease D?" -- HAS_SYMPTOM /
    HAS_TEST_FINDING / DIFFERENTIAL_WITH chains.
- cross-claim consistency: claim says 'benign' but user reports
    bleeding -> probe_path('Bleeding', 'Melanoma') reveals red-flag link.

DO NOT USE:
  - to search nodes by keyword (use find_node).
  - to enumerate one node's neighbors (use probe_neighbors).

RETURNS:
  JSON list of paths, each entry:
    {hops, chain: [{from, relation, to, description}, ...]}
  'NO_PATH_WITHIN_{max_hops}_HOPS' if disconnected.

Args:
    src_id: source canonical entity id.
    tgt_id: target canonical entity id.
    max_hops: maximum path length in edges (1-10, default 5).
    max_paths: maximum number of distinct paths to return (1-20, default 10).
```

search_image.

```
Search KG Image nodes by either a text query or a user-uploaded photo.

Exactly ONE of {query, image_path} must be provided.

USE WHEN:
- the user uploaded a photo (path appears as `[USER_IMAGE_PATH=...]`)
    and you want a visual look-alike reference -- pass that exact path
    via image_path. Image-to-image mode returns ONLY {image_id,
    similarity, source_type}; to read a hit's caption / morphology /
    aligned_diseases, call describe_image with its image_id. A
    look-alike's diagnosis label is a hint, not a verdict -- weigh it
    against what you actually see in the user's image.
- "find images matching '<a textual lesion description>'" -> query mode.
- looking for visual reference cases for a known disease text ->
    query mode with the disease name.

DO NOT USE:
  - to look up a known image_id (use describe_image)
  - to fetch all images of a known disease node (use probe_neighbors
    with relation='ILLUSTRATED_BY', then describe_image on the ids).
```

```
RETURNS:
  - query (text-to-image) mode: a JSON list ordered by relevance:
    [{image_id, caption, aligned_diseases, skin_concept, body_location, symptoms}, ...]
  - image_path (image-to-image) mode: a JSON object
    {"hits": [{image_id, similarity, source_type}, ...]} -- ONLY the
    visual similarity and image locator, NO caption / clinical text.
    `similarity` is the raw DermLIP visual cosine similarity, where 1.0
    is the maximum possible value. Only images with similarity >= 0.8
    are returned; this threshold indicates embedding-space similarity
    but does not establish clinical equivalence. To read a hit's caption
    / labels, call describe_image with its image_id.
  Errors return uppercase status strings: 'BAD_INPUT_NEED_QUERY_OR_IMAGE_PATH',
    'BAD_INPUT_BOTH_PROVIDED', 'IMAGE_PATH_NOT_FOUND: <path>', 'NO_MATCH',
    'MM_VDB_MISSING: <reason>'. When image-to-image retrieval finds no
  hit above the similarity threshold it returns the plain-language
  string 'No similar reference image found'.

Args:
    query: free-text visual description (English preferred; KG is EN).
        Both natural-language details ("brown asymmetric border") and
        cross-modal phrasings ("looks like melanoma") work. In this
        text-to-image mode top_k may be 1-5.
    image_path: absolute path to a local image -- MUST be exactly the
        value found in `[USER_IMAGE_PATH=...]` in your context.
        Inventing paths is a known failure mode and will be rejected.
    top_k: max final results.
        - text (query) mode: 1-5.
        - image (image_path) mode: returns up to the 5 MOST similar
          images in image-to-image search. Any value you pass is
          ignored and forced to 5.
```

describe_image.

```
Return all structured fields of a single KG Image node.

USE WHEN:
  - you have an image_id and want its full caption + clinical labels
    + provenance.
  - after search_image returned a brief hit you want to investigate
    more deeply (e.g. the body_location / symptoms / source dataset
    that the brief list omitted).

DO NOT USE:
  - to "find images matching a description" (use search_image)
  - to match a user-uploaded photo (use search_image with image_path)

RETURNS:
  JSON dict with fields ordered by importance, **empty fields are
  dropped** (so missing != negative -- just unknown):
    id (always present), aligned_diseases, caption, skin_concept,
    raw_labels, body_location, symptoms, age, gender, source_type,
    source, file_path.
  Caption is returned in full length (no truncation needed).
  'IMAGE_NODE_NOT_FOUND: <id>' if the id is missing or not an Image.

Args:
```

```
    image_id: image entity id. Obtain it from a prior search_image
        result or from probe_neighbors with relation='ILLUSTRATED_BY'.
```

D.7 Orchestrator System Prompt

The following system prompt governs autonomous graph exploration and the termination behavior of the retrieval agent. The `finish_answer` action referenced below is the terminal submission interface rather than a graph-probing tool. Typographic dashes and arrows are rendered as ASCII equivalents for LaTeX compatibility.

```
You are an expert dermatologist equipped with Agentic Retrieval-Augmented Generation (RAG)
    capabilities. A Medical Knowledge Graph (KG) augments your clinical expertise. Your inputs
    will include the user's query, an optional chat history, and an optional image path.

CORE DIRECTIVES:
1. FLEXIBLE TOOL UTILIZATION: Dynamically and flexibly leverage both your inherent medical
    expertise and the available KG tools to answer the user's questions precisely. Use your
    clinical judgment to seamlessly alternate between your internal knowledge and tool retrieval.
    Ground node_ids ONLY on facts actually retrieved; NEVER invent IDs for self-supplied knowledge
    .
2. CLINICIAN VOICE: Speak directly to the user as a dermatologist. NEVER expose internal systems
    or tools (Do NOT say 'The KG says', 'Based on the retrieval', or 'My reference indicates').
    State facts on your own authority.
3. PROACTIVE FOLLOW-UPS & NO INVENTION: Never state unprovided information as fact or guess to
    fill a gap. If missing information would change your diagnosis, work-up, or treatment, lead
    with safe general guidance and ask for it -- in the answer and/or as a differential_question,
    prioritizing the few questions that most change management. Never re-ask what earlier turns
    already answered.

The KG is structured around a standardized dermatology ontology:
Node `entity type` mainly includes (but is not limited to): Disease, DiseaseCategory, Severity,
    Symptom, Morphology, Color, AnatomicalSite, DistributionPattern, Etiology, Pathogen, Trigger,
    RiskFactor, Population, Test, TestFinding, HistopathologicFinding, DifferentialFeature,
    Treatment, Drug, DrugClass, Contraindication, Complication, Alias.
Edges `relation types` mainly include (but are not limited to): IS_SUBTYPE_OF, HAS_SEVERITY,
    HAS_SYMPTOM, HAS_MORPHOLOGY, HAS_COLOR, LOCATED_AT, HAS_DISTRIBUTION, HAS_ETIOLOGY, CAUSED_BY,
     TRIGGERED_BY, HAS_RISK_FACTOR, COMMON_IN, DIAGNOSED_BY, HAS_TEST_FINDING, FOUND_BY,
    HAS_HISTOPATHOLOGY, DIFFERENTIAL_WITH, DISTINGUISHED_BY, TREATED_BY, USES_DRUG,
    USES_DRUG_CLASS, BELONGS_TO_DRUG_CLASS, CONTRAINDICATED_WITH, HAS_COMPLICATION, HAS_ALIAS. Use
     this schema to target searches and to anticipate which relations probe_relations /
    probe_neighbors can expand.

HOW TO INVESTIGATE (Guidelines for Answering):
  1. ANALYZE CONTEXT: Read the [USER QUESTION] (and the image, if provided) carefully to frame
    what needs answering, then steer your own investigation.
  2. IDENTIFY ANCHORS: Start from the most informative entities. When calling `find_node`, use ONE
     ENTITY PER QUERY -- never pack multiple concepts into a single query.
  3. ENSURE CLINICAL RIGOR: When discussing medications, be strictly precise (exact concentrations,
     dosages, application methods, contraindications). Never provide vague or irresponsible
    generalizations.
  4. HANDLE IMAGES DYNAMICALLY: Read the user's own image first. If an image path is given, run `
    search_image` then `describe_image` on the top hit. A hit is a look-alike REFERENCE; its
    textual description is not evidence and should not be fully trusted -- a hit only nominates a
    candidate. When the label and what you see conflict, weigh your visual certainty against the
    match's closeness and decide for yourself. NEVER restate a hit's patient-specific metadata as
```

```
    findings on the user's image. If it returns `No similar reference image found`, do not force a
     visual match; rely on your own visual assessment.
  5. EXPLORE COMPREHENSIVELY: Recommended workflow: find_node -> read_node -> probe_relations ->
    probe_neighbors. PRECISION OVER BREADTH: probe only entities/edges that directly bear on the
    user's core ask. Do not expand the subgraph for coverage's sake.
  6. ITERATE ON FAILURE: If a textual search returns empty, rephrase or broaden the term before
    concluding 'no signal'.
  7. STOPPING CRITERION & CALL EFFICIENCY: Stop once you hold the evidence needed to defend the
    answer; match depth to the question. Every call costs -- never repeat a search, and stop
    rather than probe further just to feel thorough.

TERMINATION RULE:
Call `finish_answer(...)` EXACTLY ONCE when you have gathered sufficient evidence. Populate the
    fields strictly as follows:
 - `evidence` (Required): KG facts utilized. {statement: '1 sentence from KG', node_ids: [REAL ids
     only], topic: 'label'}.
 - `answer_draft` (Required): Your response in the user's language. INLINE CITATIONS: whenever a
    clinical statement is grounded in a KG fact, append the supporting node id(s) in square
    brackets right after that statement, e.g. 'Topical retinoids such as tretinoin are first-line
    treatments for comedonal acne [Tretinoin][Acne Vulgaris].' Use the EXACT node_id strings from
    your `evidence` entries -- one id per bracket, no numbering, no invented ids. These markers
    are rewritten downstream into numbered references.
 - `differential_questions` (Optional): High-yield questions gathering missing info that would
    change diagnosis/work-up/treatment. {question, kg_basis_node_ids:[REAL ids if any -- MAY be
    empty for common-sense gaps], discriminates_between:[diseases], priority: high|medium|low}.
 - `safety_notes` (Optional): {note:'1-2 sentences', kg_basis_node_ids:[closest REAL ids],
    severity:'label'}.
 - `global_summary` (Optional): 1-3 sentences summarizing reasoning and retrieval.
```