arXiv:2611.11010v1 [cs.AI] 7 Nov 2026

LIGHTFOOT: Graph Retrieval-Augmented Generation with Exact, Model-Free Document Deletion

Uday Allu

Abstract

Graph-based retrieval-augmented generation (GraphRAG) improves multi-hop question answering, but deployed systems face two costs that research comparisons rarely measure: building the graph, and removing documents from it. We present LIGHTFOOT, a GraphRAG system in which every entity, link, fact and synonym edge is a deterministic function of per-passage extractions carrying chunk-level provenance. Indexing makes one language-model call per passage, retrieval makes none, and deleting a document is a database operation with no model calls. On the full 11,402-passage MuSiQue corpus we verify that, after deleting 10% of the documents, the graph is identical to a fresh index of the surviving documents. We compare LIGHTFOOT with HippoRAG 2, LightRAG and fast-graphrag under identical extraction, answer and embedding models on MuSiQue, 2WikiMultiHopQA and HotpotQA. LIGHTFOOT builds the cheapest index on all three datasets, 26–31% below HippoRAG 2 and 3–6× below LightRAG and fast-graphrag. Against HippoRAG 2 it is significantly better on MuSiQue retrieval (+2.8 points recall@5 and +4.3 recall@10 on a held-out split) and on HotpotQA answers (+7.7 points exact match), and significantly worse on 2WikiMultiHopQA (−3.9 recall@2, −8.4 F1). Deleting the same documents took LightRAG 2,187 model calls and 3.7 hours; fast-graphrag has no deletion path. Code and per-question results: https://github.com/udayallu-ml/Lightfoot-GraphRAG.

1 Introduction

Retrieval-augmented generation [14] grounds a language model in documents retrieved at query time. Plain similarity search fails on multi-hop questions whose evidence is spread over passages that the question never names. GraphRAG systems address this by extracting entities and relations into a graph and traversing it at query time [4, 8–10]; surveys catalogue dozens of variants [11, 20].

Published comparisons focus on answer accuracy. Two properties that decide whether a system can be operated in an organisation receive little attention. The first is *index cost*: extraction dominates the bill, and frameworks differ by an order of magnitude because of second extraction passes, model-written entity summaries and community reports. The second is *deletion*. Retention policies, contract expiry and the right to erasure [6] require removing a document and everything derived from it. Summaries that blend information from many documents make this expensive or impossible. A merged entity description must be regenerated, and a community report cannot be attributed to its sources at all. Machine unlearning studies the analogous problem for model weights [1]. For retrieval indexes the problem is usually unmeasured.

We surveyed the source code of 17 open-source GraphRAG frameworks. Five offer document-level deletion, and none claims that the result equals a graph that never contained the document.

Contributions.

1. A provenance model in which every derived record is a pure function of the surviving per-passage extractions (§3.2). Deletion is therefore exact and needs no model calls. We verify this at full corpus scale.
2. A one-call extraction schema and a retrieval pipeline with no model calls: fact-seeded Personalized PageRank fused with a local cross-encoder (§3.3).
3. A cost-normalised, deletion-audited comparison against three open-source baselines under identical models and prompts, with paired significance tests (§5).
4. Negative and diagnostic results: a model-based entity-merge judge, a hybrid keyword and embedding search channel, and larger answer contexts gave no reliable gain. We also document two silent baseline failure modes uncovered during benchmarking (§6).

2 Related work

Community-summary GraphRAG. Microsoft GraphRAG [4] clusters an extracted entity graph and asks a model to write a report for each community. This supports corpus-level questions but makes indexing the most expensive in our survey and ties every report to many source documents.

Keyword dual-level retrieval. LightRAG [8] extracts entities and relations with model-written descriptions, merges duplicate descriptions with further model calls, and retrieves by matching model-extracted keywords against entity and relation vectors.

PageRank over extracted facts. HippoRAG [9] and HippoRAG 2 [10] seed Personalized PageRank [12] from entities linked to the query and rank passages by the resulting mass. HippoRAG 2 matches the query against fact embeddings and filters candidate facts with a model call. fast-graphrag [2] also ranks with PageRank, after one structured extraction call per chunk.

Other designs. RAPTOR [16] builds a tree of recursive summaries. KAG [15] plans retrieval over a schema-constrained graph with logical forms. LIGHTFOOT is closest to the PageRank family but differs in three ways: it makes no model call during retrieval, it fuses a cross-encoder with the walk, and its derived state is provenance-exact.

3 Method

3.1 Extraction

Each passage is processed by exactly one call to an open 70B instruction model [7]. The call returns minified JSON with up to 12 entities, each with a name, a type from a configurable list and a description of at most 8 words. It also returns up to 12 facts, each written as an array of subject, predicate, object and a verbatim evidence quote of at most 12 words. Writing facts as arrays and forbidding pretty-printing reduced output from 436 to 269 tokens per passage on a 32-passage sample. Over the same sample, facts per passage rose from 3.8 to 7.7. On MuSiQue the extractor averaged 442 input and 310 output tokens per passage. Extractions are cached under a key that includes a hash of the prompt and schema.

3.2 Provenance model and the erasure invariant

Each occurrence of a name in a chunk is stored as a *form* (c, n, t): chunk c, normalised name n, type t, plus that chunk's description of the entity. Two forms of the same type are *linked* when the

cosine similarity of their name embeddings is at least $\tau_m = 0.92$. Name embeddings are computed in fixed-shape batches, so a name's vector does not depend on which other names share its batch. An *entity* is a connected component of live forms. Its identifier is derived from its lexicographically smallest member name, and its canonical name is the alias with the most provenance, with ties broken alphabetically. *Synonym edges* join every pair of same-type entities whose name similarity is at least $\tau_s = 0.80$, without a top-k cap. *Facts* are keyed by their subject and object forms and carry the set of chunks that asserted them.

Proposition 1. *Let $G(D)$ denote the store produced by indexing document set D. For any partition $D = A \cup B$, deleting B from $G(D)$ yields a store equal to $G(A)$.*

Proof sketch. Every stored record is a function of the set of live forms and their chunks. Forms and fact assertions are per-chunk and are removed with their chunk. Form links depend only on the two forms involved. Components therefore depend only on the set of live forms, never on insertion order, batching or history, and so do entity identifiers and canonical names. Synonym edges are a threshold function of the live entity names, with no top-k lists that could depend on removed neighbours. No record is produced by a model call that reads several chunks. Deleting B removes exactly the forms and assertions contributed by B's chunks, then recomputes the components and synonym edges it touched. That recomputation yields the same records as computing them from A alone. □

Two design choices follow directly from the invariant. A model-based judge for ambiguous name pairs would make links depend on what the judge saw; we removed it after measuring that it changed no retrieval metric beyond noise (§6). A top-k synonym list would depend on deleted neighbours; it is replaced by a pure threshold.

3.3 Retrieval

The graph has entity and passage nodes. Passage–entity edges come from forms, entity–entity edges are weighted by the number of facts linking the pair, and synonym edges are weighted by similarity. Given a question q:

1. Embed q with bge-base-en-v1.5 [18].
2. Seed a Personalized PageRank walk. The subject and object entities of the two facts most similar to q receive mass equal to the fact similarity divided by the number of passages mentioning the entity. The 100 passages most similar to q receive mass up to 0.05 after min–max normalisation.
3. Run the walk with damping 0.9 and take the 30 highest-scoring passages.
4. Score these passages with a local cross-encoder (bge-reranker-base) and fuse the cross-encoder rank with the PageRank rank by reciprocal rank fusion with $k = 60$ [3].

The fused ranking is the system's retrieval output. The answer model receives the ten most similar facts with their evidence quotes and the top 10 passages in ranked order, and returns a short answer span as JSON. No step before the answer call uses a language model.

4 Experimental setup

Datasets. We use MuSiQue [17], 2WikiMultiHopQA [13] and HotpotQA [19], following HippoRAG's construction: a seeded sample of development questions, with the corpus formed by all supporting and distractor paragraphs of the sampled questions. MuSiQue has 1,000 questions over 11,402 passages, 2WikiMultiHopQA 300 over 2,131, and HotpotQA 300 over 2,957. Retrieval is scored on every question. Answers are generated for the first 300 questions of each dataset.

Table 1: Retrieval recall@5 and answer accuracy. MuSiQue answers are scored on the 232 of the first 300 questions whose evidence was not deleted; 2WikiMultiHopQA and HotpotQA answers on all 300. Best value per column in bold.

	MuSiQue			2WikiMultiHopQA			HotpotQA		
System	R@5	EM	F1	R@5	EM	F1	R@5	EM	F1
LIGHTFOOT	**.646**	**.401**	**.531**	.884	.560	.633	**.943**	**.630**	**.783**
HippoRAG 2	.616	.366	.511	**.905**	**.600**	**.717**	.930	.553	.734
LightRAG	.480	.306	.440	.722	.510	.581	.902	.543	.709
fast-graphrag	.412	.237	.332	.838	.523	.582	.732	.517	.671
Dense retrieval only	.544	–	–	–	–	–	–	–	–

Baselines and parity. We compare against HippoRAG 2, LightRAG and fast-graphrag, which are the strongest maintained systems in their families that produce a ranked passage list. Every system uses Llama 3.3 70B for extraction, gpt-5.6-luna with reasoning disabled for answers, and bge-base-en-v1.5 for embeddings. Every system receives the same instruction to return the shortest answer span and always give a best guess. Without this instruction, two baselines answered "no relevant information" or "cannot be determined" on a large share of questions. We report scores with the shared instruction. LightRAG runs in its default `mix` mode on PostgreSQL storage, because its default file storage rewrites every graph and vector file for each inserted document (§6).

Tuning protocol. LIGHTFOOT's retrieval settings were selected on MuSiQue questions 1–500 only. MuSiQue comparisons with HippoRAG 2 use questions 501–1000. 2WikiMultiHopQA and HotpotQA were never used for tuning.

Metrics. We report recall@k over gold supporting passages, and exact match (EM) and token F1 over answer spans with standard normalisation. Index cost counts every model call and prices its tokens at one fixed list price per model: \$0.10/\$0.32 per million input/output tokens for Llama 3.3 70B and \$0.20/\$1.20 for gpt-5.6-luna. This keeps provider routing from distorting the comparison. Significance uses paired bootstrap 95% intervals over questions [5]. A difference is called significant when its interval excludes zero.

Deletion audit. After the main query pass, a seeded random 10% of documents is deleted through each system's own API, and every question is re-run. Recall is reported separately for questions whose evidence survived and for questions whose evidence was deleted. For LIGHTFOOT we additionally fingerprint the post-deletion store and compare it with a fresh index of the surviving documents, built in a separate process from the extraction cache.

Hardware. All local computation, including embeddings, reranking and graph walks, ran on one Apple M3 Pro laptop with 18 GB of memory. Model calls went through OpenRouter.

5 Results

Accuracy. Table 1 reports recall@5 and answer accuracy, and Table 2 the paired comparison with the strongest baseline. Of the 15 comparisons with HippoRAG 2, LIGHTFOOT is significantly better on 5, significantly worse on 2 (both on 2WikiMultiHopQA) and tied on 8. It exceeds LightRAG

Table 2: LIGHTFOOT minus HippoRAG 2, paired by question, in percentage points with 95% bootstrap intervals. MuSiQue retrieval uses the held-out questions 501–1000.

Metric	MuSiQue	2WikiMultiHopQA	HotpotQA
Recall@2	−0.4 [−2.8, +2.0]	**−3.9** [−7.0, −0.8]	+1.8 [−1.5, +5.2]
Recall@5	**+2.8** [+0.5, +5.1]	−2.1 [−4.5, +0.3]	+1.3 [−0.7, +3.2]
Recall@10	**+4.3** [+2.5, +6.2]	+0.1 [−1.4, +1.7]	**+1.3** [+0.2, +2.7]
EM	+3.4 [−2.2, +9.1]	−4.0 [−9.0, +1.0]	**+7.7** [+3.0, +12.3]
F1	+2.0 [−3.4, +7.4]	**−8.4** [−12.9, −4.2]	**+5.0** [+1.3, +8.7]

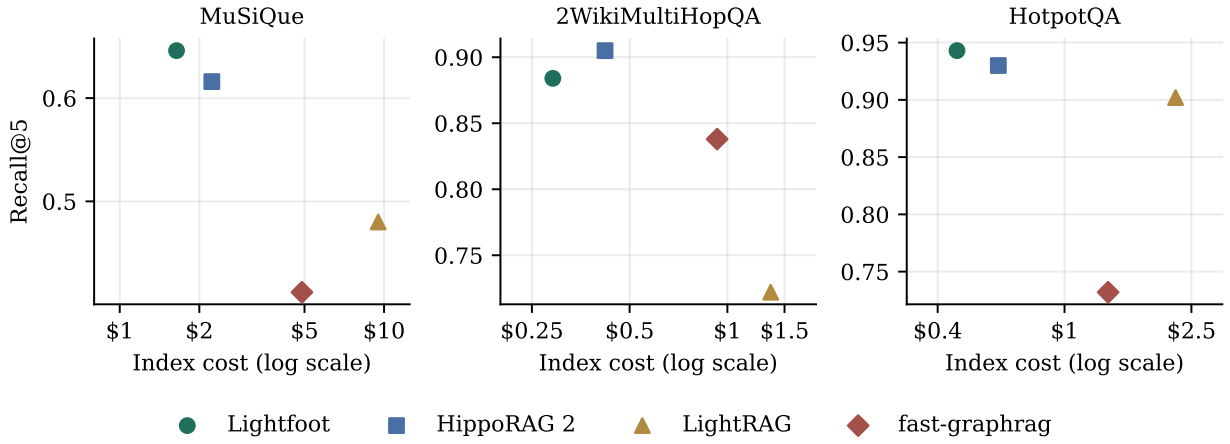

Figure 1: Recall@5 against normalised index cost for each dataset. Points nearer the top left are better.

and fast-graphrag at recall@5 on every dataset, and exceeds dense retrieval alone by 10.2 points recall@5 on MuSiQue.

Cost. Table 3 and Figure 1 show index cost. LIGHTFOOT is cheapest on every dataset: 26–31% below HippoRAG 2, which makes two extraction calls per passage, and 3–6× below LightRAG and fast-graphrag. At query time, retrieval in LIGHTFOOT makes no model call. Each baseline makes one per question, for keyword extraction (LightRAG), fact filtering (HippoRAG 2) or query entity extraction (fast-graphrag). Answer cost is similar across systems, about $0.06 per 100 answers for LIGHTFOOT.

Latency. Median retrieval-only latency on MuSiQue, excluding the answer call, was 0.53 s for LIGHTFOOT, 0.81 s for LightRAG, 1.56 s for HippoRAG 2 and 1.96 s for fast-graphrag. Latency including the answer call varied with provider load between runs and is not compared.

Deletion. Table 4 reports the deletion audit on MuSiQue. For every system that supports deletion, recall on questions whose evidence survived did not decrease. Recall on questions whose evidence was deleted fell sharply, which confirms that the information left retrieval. LightRAG re-summarises the descriptions of every entity touched by a deleted document, which took 2,187 model calls and 3.7 hours. On 2WikiMultiHopQA and HotpotQA it took 165 and 330 calls. LIGHTFOOT and HippoRAG 2 made no model calls.

Table 3: Index cost at normalised list prices, and model calls per passage. LightRAG's MuSiQue cost is estimated from its 28,529 extraction and summary calls at measured average tokens per call, because its extraction was billed across several runs.

System	MuSiQue	2WikiMultiHopQA	HotpotQA	Calls per passage
LIGHTFOOT	**$1.64**	**$0.29**	**$0.46**	1.0
HippoRAG 2	$2.23	$0.42	$0.62	2.0
fast-graphrag	$4.89	$0.93	$1.37	1.1
LightRAG	≈$9.5	$1.36	$2.23	2.5

Table 4: Deleting a seeded 10% of MuSiQue documents (1,140 of 11,402). Recall@5 on questions whose evidence survived (intact) and on questions whose evidence was deleted.

System	Time	Model calls	Intact, before	Intact, after	Deleted, after
LIGHTFOOT	74 s	0	.648	.661	.369
HippoRAG 2	9 s	0	.625	.639	.333
LightRAG	13,294 s	2,187	.491	.507	.292
fast-graphrag	no deletion API; rebuilding from scratch is required				

Exactness at full scale. After the deletion, the LIGHTFOOT store was fingerprinted in one process. A fresh index of the 10,262 surviving documents was built in a second process from the extraction cache. The fingerprints were identical, and every table matched exactly: 68,455 entities, 71,195 aliases, 106,472 mentions, 92,608 facts, 94,032 fact sources and 63,514 synonym edges. The fresh build made no model calls. An earlier version that resolved new names against the current canonical entities, and kept top-k synonym lists, differed from the fresh index by 12 entities and 340 synonym edges. That gap is what motivated the pairwise, history-free design of §3.2.

6 Analysis

Ablations. All ablations were run on MuSiQue questions 1–500 without model calls. Fusing the cross-encoder rank with the PageRank rank, instead of replacing the PageRank order, was the largest single gain at 3.4 points of recall@5. Raising damping from 0.5 to 0.9 under fusion added 2 to 3 points, and seeding from 2 facts rather than 5 added about 1. Removing fact nodes from the graph, by turning facts into weighted entity–entity edges, changed nothing and shrank the graph by 40%. On a 2,000-passage slice, removing the model-based judge for name pairs with similarity 0.80 to 0.92 changed recall@5 by less than one point, within noise across 141 questions. It saved 5,626 judge calls on the full corpus.

Answer context. Replacing eight passages plus entity descriptions with the top 10 passages in ranked order raised HotpotQA EM by 2.3 points, with a 95% interval of +0.3 to +4.7. Changes on the other two datasets were within noise. Adding a separate hybrid channel, Elasticsearch BM25 and dense search fused by reciprocal rank fusion, raised the share of gold passages in the context from 0.852 to 0.873 on a MuSiQue slice. It added 35% to answer tokens and produced no reliable answer gain. Every way of selecting 10 passages from the union of the graph and hybrid rankings covered fewer gold passages than the graph ranking alone.

Where Lightfoot fails. On the 300 answered MuSiQue questions, in the run that also included the hybrid channel, 183 answers were wrong. In 109 of these at least one gold passage was missing

from the context, and 74 had every gold passage present. Missing evidence concentrates on later hops. The entity of the first hop's passage is named in the question 83% of the time, against 36% for the second hop and 28% for the third and fourth. Search over the question text, lexical or dense, cannot reach these passages; only following the chain can. Seeding from the two most similar facts sometimes selects incidental facts such as release dates over the bridging fact. HippoRAG 2's model-based fact filter addresses exactly this, which plausibly explains its advantage on 2WikiMultiHopQA, where most questions compare two named entities. Of the 74 failures with complete evidence, 49 were near misses with F1 of at least 0.5: over-long spans, answers naming an intermediate entity, and off-by-one dates.

Answer noise. Between two runs whose contexts differed only in passage count (8 vs. 10), about 30 of 300 answers flipped between correct and incorrect while the net change was near zero. Answer differences below roughly 4 points on 300 questions should therefore be read as ties, which is how Table 2 treats them.

Baseline failure modes. Two silent failures would have corrupted the comparison had they gone unnoticed. LightRAG's first index aborted most of its insertion pipeline after an embedding timeout but returned normally, leaving 31% of chunks indexed. Its first score (recall@5 0.239) was discarded, and the reported run verifies that every document reached the processed state. LightRAG's default file storage rewrote all graph and vector files for each inserted document, about 415 MB per document at 3,000 passages, so throughput fell continuously. HippoRAG 2 aborted its entire index when 24 of 11,402 extraction outputs were malformed JSON. A tolerant parser in our adapter recovered 11,397; five passages were indexed without entities.

7 Limitations

- LIGHTFOOT's retrieval settings were tuned on half of one dataset. The held-out half and the two untuned datasets agree on cost, and accuracy is mixed: one untuned dataset is a win and the other a loss.
- Answer metrics rest on 232 to 300 questions per dataset, and 2WikiMultiHopQA and HotpotQA use 300-question samples rather than the 1,000 common in prior work.
- All systems share one extraction model and one embedder. HippoRAG 2 is therefore evaluated without the 7B embedder used in its paper, and its absolute scores are below those published.
- Deletion exactness was verified for LIGHTFOOT only. HippoRAG 2's provenance-based deletion may also be exact.
- Corpus-level summarisation questions, the target of community-summary systems, were not evaluated.

8 Conclusion

LIGHTFOOT shows that a GraphRAG index can be cheap to build, need no model at retrieval time, and support exact, model-free deletion, while matching the strongest open-source baseline on multi-hop accuracy overall. The remaining accuracy gap lies in selecting seed facts and reaching later hops that the question never names. Model-free second-hop retrieval from entities found in the first round is the most direct next step.

Reproducibility. Code, the benchmark harness, baseline adapters, run configurations, the question and passage subsets, and per-question predictions and cost ledgers for every reported run are available at https://github.com/udayallu-ml/Lightfoot-GraphRAG.

Use of AI tools. An AI assistant was used to implement parts of the benchmarking harness, to run experiments and to draft this manuscript. The authors reviewed all results and text and take full responsibility for them.

References

[1] Lucas Bourtoule, Varun Chandrasekaran, Christopher A. Choquette-Choo, Hengrui Jia, Adelin Travers, Baiwu Zhang, David Lie, and Nicolas Papernot. Machine unlearning. In *IEEE Symposium on Security and Privacy*, 2021.

[2] Circlemind. fast-graphrag. https://github.com/circlemind-ai/fast-graphrag, 2024.

[3] Gordon V. Cormack, Charles L. A. Clarke, and Stefan Buettcher. Reciprocal rank fusion outperforms Condorcet and individual rank learning methods. In *Proceedings of the 32nd International ACM SIGIR Conference on Research and Development in Information Retrieval*, 2009.

[4] Darren Edge, Ha Trinh, Newman Cheng, Joshua Bradley, Alex Chao, Apurva Mody, Steven Truitt, and Jonathan Larson. From local to global: A graph RAG approach to query-focused summarization. *arXiv preprint arXiv:2404.16130*, 2024.

[5] Bradley Efron and Robert J. Tibshirani. *An Introduction to the Bootstrap*. Chapman and Hall/CRC, 1994.

[6] European Union. Regulation (EU) 2016/679 of the European Parliament and of the Council (General Data Protection Regulation), article 17: Right to erasure. Official Journal of the European Union, L 119, 2016.

[7] Aaron Grattafiori et al. The Llama 3 herd of models. *arXiv preprint arXiv:2407.21783*, 2024.

[8] Zirui Guo, Lianghao Xia, Yanhua Yu, Tu Ao, and Chao Huang. LightRAG: Simple and fast retrieval-augmented generation. *arXiv preprint arXiv:2410.05779*, 2024.

[9] Bernal Jiménez Gutiérrez, Yiheng Shu, Yu Gu, Michihiro Yasunaga, and Yu Su. HippoRAG: Neurobiologically inspired long-term memory for large language models. In *Advances in Neural Information Processing Systems*, 2024.

[10] Bernal Jiménez Gutiérrez, Yiheng Shu, Weijian Qi, Sizhe Zhou, and Yu Su. From RAG to memory: Non-parametric continual learning for large language models. *arXiv preprint arXiv:2502.14802*, 2025.

[11] Haoyu Han et al. Retrieval-augmented generation with graphs (GraphRAG). *arXiv preprint arXiv:2501.00309*, 2025.

[12] Taher H. Haveliwala. Topic-sensitive PageRank. In *Proceedings of the 11th International World Wide Web Conference*, 2002.

[13] Xanh Ho, Anh-Khoa Duong Nguyen, Saku Sugawara, and Akiko Aizawa. Constructing a multi-hop QA dataset for comprehensive evaluation of reasoning steps. In *Proceedings of the 28th International Conference on Computational Linguistics*, 2020.

[14] Patrick Lewis, Ethan Perez, Aleksandra Piktus, Fabio Petroni, Vladimir Karpukhin, Naman Goyal, Heinrich Küttler, Mike Lewis, Wen-tau Yih, Tim Rocktäschel, Sebastian Riedel, and Douwe Kiela. Retrieval-augmented generation for knowledge-intensive NLP tasks. In *Advances in Neural Information Processing Systems*, 2020.

[15] Lei Liang et al. KAG: Boosting LLMs in professional domains via knowledge augmented generation. *arXiv preprint arXiv:2409.13731*, 2024.

[16] Parth Sarthi, Salman Abdullah, Aditi Tuli, Shubh Khanna, Anna Goldie, and Christopher D. Manning. RAPTOR: Recursive abstractive processing for tree-organized retrieval. In *International Conference on Learning Representations*, 2024.

[17] Harsh Trivedi, Niranjan Balasubramanian, Tushar Khot, and Ashish Sabharwal. MuSiQue: Multihop questions via single-hop question composition. *Transactions of the Association for Computational Linguistics*, 10:539–554, 2022.

[18] Shitao Xiao, Zheng Liu, Peitian Zhang, and Niklas Muennighoff. C-pack: Packaged resources to advance general Chinese embedding. *arXiv preprint arXiv:2309.07597*, 2023.

[19] Zhilin Yang, Peng Qi, Saizheng Zhang, Yoshua Bengio, William W. Cohen, Ruslan Salakhutdinov, and Christopher D. Manning. HotpotQA: A dataset for diverse, explainable multi-hop question answering. In *Proceedings of the 2018 Conference on Empirical Methods in Natural Language Processing*, 2018.

[20] Qinggang Zhang et al. A survey of graph retrieval-augmented generation for customized large language models. *arXiv preprint arXiv:2501.13958*, 2025.